

# Capstone Project (SEPHORA)

## Enhancing Customer Experience: Unraveling Sephora’s Skincare Reviews

### Product Dataset Overview

Feature	Description
product_id	The unique identifier for the product from the site
product_name	The full name of the product
brand_id	The unique identifier for the product brand from the site
brand_name	The full name of the product brand
loves_count	The number of people who have marked this product as a favorite
rating	The average rating of the product based on user reviews
reviews	The number of user reviews for the product
size	The size of the product, which may be in oz, ml, g, packs, or other units depending on the product type
variation_type	The type of variation parameter for the product (e.g. Size, Color)
variation_value	The specific value of the variation parameter for the product (e.g. 100 mL, Golden Sand)
variation_desc	A description of the variation parameter for the product (e.g. tone for fairest skin)
ingredients	A list of ingredients included in the product, for example: ['Product variation 1:', 'Water, Glycerin', 'Product variation 2:', 'Talc, Mica'] or if no variations ['Water, Glycerin']
price_usd	The price of the product in US dollars
value_price_usd	The potential cost savings of the product, presented on the site next to the regular price
sale_price_usd	The sale price of the product in US dollars
limited_edition	Indicates whether the product is a limited edition or not (1-true, 0-false)
new	Indicates whether the product is new or not (1-true, 0-false)
online_only	Indicates whether the product is only sold online or not (1-true, 0-false)
out_of_stock	Indicates whether the product is currently out of stock or not (1 if true, 0 if false)
sephora_exclusive	Indicates whether the product is exclusive to Sephora or not (1 if true, 0 if false)
highlights	A list of tags or features that highlight the product's attributes (e.g. ['Vegan', 'Matte Finish'])
primary_category	First category in the breadcrumb section
secondary_category	Second category in the breadcrumb section
tertiary_category	Third category in the breadcrumb section
child_count	The number of variations of the product available
child_max_price	The highest price among the variations of the product
child_min_price	The lowest price among the variations of the product

This table organizes the various features and their descriptions in a clear tabular format. Each row represents a specific feature, and the corresponding descriptions are provided in the adjacent cell.

```
In [37]: # importing necessary libraries

import pandas as pd
import numpy as np

# loading in products data

data = pd.read_csv(r"C:\Users\logeshwar\OneDrive\Documents\Sephora\product_info.csv")
data.info(verbose=True)
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8494 entries, 0 to 8493
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   product_id                            8494 non-null   object
1   product_name                          8494 non-null   object
2   brand_id                              8494 non-null   int64
3   brand_name                            8494 non-null   object
4   loves_count                           8494 non-null   int64
5   rating                                8216 non-null   float64
6   reviews                              8216 non-null   float64
7   size                                  6863 non-null   object
8   variation_type                        7050 non-null   object
9   variation_value                       6896 non-null   object
10  variation_desc                        1250 non-null   object
11  ingredients                           7549 non-null   object
12  price_usd                             8494 non-null   float64
13  value_price_usd                       451 non-null    float64
14  sale_price_usd                        270 non-null    float64
15  limited_edition                       8494 non-null   int64
16  new                                    8494 non-null   int64
17  online_only                           8494 non-null   int64
18  out_of_stock                          8494 non-null   int64
19  sephora_exclusive                     8494 non-null   int64
20  highlights                            6287 non-null   object
21  primary_category                      8494 non-null   object
22  secondary_category                    8486 non-null   object
23  tertiary_category                     7504 non-null   object
24  child_count                           8494 non-null   int64
25  child_max_price                       2754 non-null   float64
26  child_min_price                       2754 non-null   float64
dtypes: float64(7), int64(8), object(12)
memory usage: 1.7+ MB

```

From this, we gather that the products dataset contains 27 columns and 8,494 rows of data. The columns vary in data types, and there is missing data within the dataframe.

## Review Dataset Overview

Feature	Description
author_id	The unique identifier for the author of the review on the website
rating	The rating given by the author for the product on a scale of 1 to 5
is_recommended	Indicates if the author recommends the product or not (1-true, 0-false)
helpfulness	The ratio of all ratings to positive ratings for the review: $\text{helpfulness} = \frac{\text{total\_pos\_feedback\_count}}{\text{total\_feedback\_count}}$
total_feedback_count	Total number of feedback (positive and negative ratings) left by users for the review
total_neg_feedback_count	The number of users who gave a negative rating for the review
total_pos_feedback_count	The number of users who gave a positive rating for the review
submission_time	Date the review was posted on the website in the 'yyyy-mm-dd' format
review_text	The main text of the review written by the author
review_title	The title of the review written by the author
skin_tone	Author's skin tone (e.g. fair, tan, etc.)
eye_color	Author's eye color (e.g. brown, green, etc.)
skin_type	Author's skin type (e.g. combination, oily, etc.)
hair_color	Author's hair color (e.g. brown, auburn, etc.)
product_id	The unique identifier for the product on the website

This table presents the various features related to the review data in a tabular format. Each row corresponds to a specific feature, and the descriptions are provided in the adjacent cell, making it easy to read and understand the details of each feature.

For the datasets related to the reviews, it will be helpful to join them all into one dataframe, which we can do with the following:

```

In [38]: # getting the files

import pandas as pd

t1 = pd.read_csv(r"C:\Users\logeshwar\OneDrive\Documents\Sephora\reviews_0_250.csv", low_memory=False)
t2 = pd.read_csv(r"C:\Users\logeshwar\OneDrive\Documents\Sephora\reviews_250_500.csv", low_memory=False)
t3 = pd.read_csv(r"C:\Users\logeshwar\OneDrive\Documents\Sephora\reviews_500_750.csv", low_memory=False)
t4 = pd.read_csv(r"C:\Users\logeshwar\OneDrive\Documents\Sephora\reviews_750_1000.csv", low_memory=False)
t5 = pd.read_csv(r"C:\Users\logeshwar\OneDrive\Documents\Sephora\reviews_1000_1500.csv", low_memory=False)
t6 = pd.read_csv(r"C:\Users\logeshwar\OneDrive\Documents\Sephora\reviews_1500_end.csv", low_memory=False)

# combining the dfs

```

```
texta = pd.concat([t1,t2,t3,t4,t5,t6])
texta.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1301136 entries, 0 to 49976
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            1301136 non-null  int64
1   author_id                             1301136 non-null  object
2   rating                                1301136 non-null  int64
3   is_recommended                        1107162 non-null  float64
4   helpfulness                           631670 non-null  float64
5   total_feedback_count                  1301136 non-null  int64
6   total_neg_feedback_count              1301136 non-null  int64
7   total_pos_feedback_count              1301136 non-null  int64
8   submission_time                       1301136 non-null  object
9   review_text                           1299520 non-null  object
10  review_title                           930754 non-null  object
11  skin_tone                              1103798 non-null  object
12  eye_color                              1057734 non-null  object
13  skin_type                              1172830 non-null  object
14  hair_color                             1037824 non-null  object
15  product_id                             1301136 non-null  object
16  product_name                           1301136 non-null  object
17  brand_name                             1301136 non-null  object
18  price_usd                              1301136 non-null  float64
dtypes: float64(3), int64(5), object(11)
memory usage: 198.5+ MB
```

We gather that there are around 1.3 million reviews with varying amounts of missing data across 19 columns of the 6 datasets.

It can also be helpful to view the numeric and non-numeric columns of both dataframes. We can do so with the following:

#### Numeric Columns

```
In [39]: numeric_cols = data.select_dtypes(include = ['number']).columns
print(numeric_cols)
print(f'{len(numeric_cols)} Numeric Columns in Products Dataset')

Index(['brand_id', 'loves_count', 'rating', 'reviews', 'price_usd',
       'value_price_usd', 'sale_price_usd', 'limited_edition', 'new',
       'online_only', 'out_of_stock', 'sephora_exclusive', 'child_count',
       'child_max_price', 'child_min_price'],
      dtype='object')
15 Numeric Columns in Products Dataset
```

```
In [40]: numeric_cols_reviews = texta.select_dtypes(include = ['number']).columns
print(numeric_cols_reviews)
print(f'{len(numeric_cols_reviews)} Numeric Columns in Reviews Dataset')

Index(['Unnamed: 0', 'rating', 'is_recommended', 'helpfulness',
       'total_feedback_count', 'total_neg_feedback_count',
       'total_pos_feedback_count', 'price_usd'],
      dtype='object')
8 Numeric Columns in Reviews Dataset
```

#### Non-Numeric Columns

```
In [41]: non_numeric_cols = data.select_dtypes(exclude=['number']).columns
print(non_numeric_cols)
print(f'{len(non_numeric_cols)} Non-Numeric Columns in Products Dataset')

Index(['product_id', 'product_name', 'brand_name', 'size', 'variation_type',
       'variation_value', 'variation_desc', 'ingredients', 'highlights',
       'primary_category', 'secondary_category', 'tertiary_category'],
      dtype='object')
12 Non-Numeric Columns in Products Dataset
```

```
In [42]: non_numeric_rev_cols = texta.select_dtypes(exclude=['number']).columns
print(non_numeric_rev_cols)
print(f'{len(non_numeric_rev_cols)} Non-Numeric Columns in Reviews Dataset')

Index(['author_id', 'submission_time', 'review_text', 'review_title',
       'skin_tone', 'eye_color', 'skin_type', 'hair_color', 'product_id',
       'product_name', 'brand_name'],
      dtype='object')
11 Non-Numeric Columns in Reviews Dataset
```

## Data Cleaning

### Data Cleaning on product Dataset

Next, we need to clean the products dataset by performing the following tasks:

Assessing missing data and determining its extent. Removing unnecessary data that is not relevant to our analysis. Identifying and handling outliers, if present. Reformatting the data if required.


Here we use the following code to see the amount of nulls by column within the dataset:

```
In [43]: num_missing = data.isna().sum()  
num_missing
```

```
Out[43]: product_id          0  
product_name          0  
brand_id              0  
brand_name            0  
loves_count           0  
rating                278  
reviews               278  
size                  1631  
variation_type        1444  
variation_value       1598  
variation_desc        7244  
ingredients            945  
price_usd              0  
value_price_usd       8043  
sale_price_usd        8224  
limited_edition         0  
new                    0  
online_only            0  
out_of_stock           0  
sephora_exclusive      0  
highlights            2207  
primary_category       0  
secondary_category     8  
tertiary_category      990  
child_count            0  
child_max_price        5740  
child_min_price        5740  
dtype: int64
```

```
In [44]: # heatmap to visualize missing data (products)  
  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(10,8))  
  
cols= data.columns  
colors=['navy','grey']  
sns.heatmap(data[cols].isna(),cmap=sns.color_palette(colors))
```

```
Out[44]: <AxesSubplot:>
```



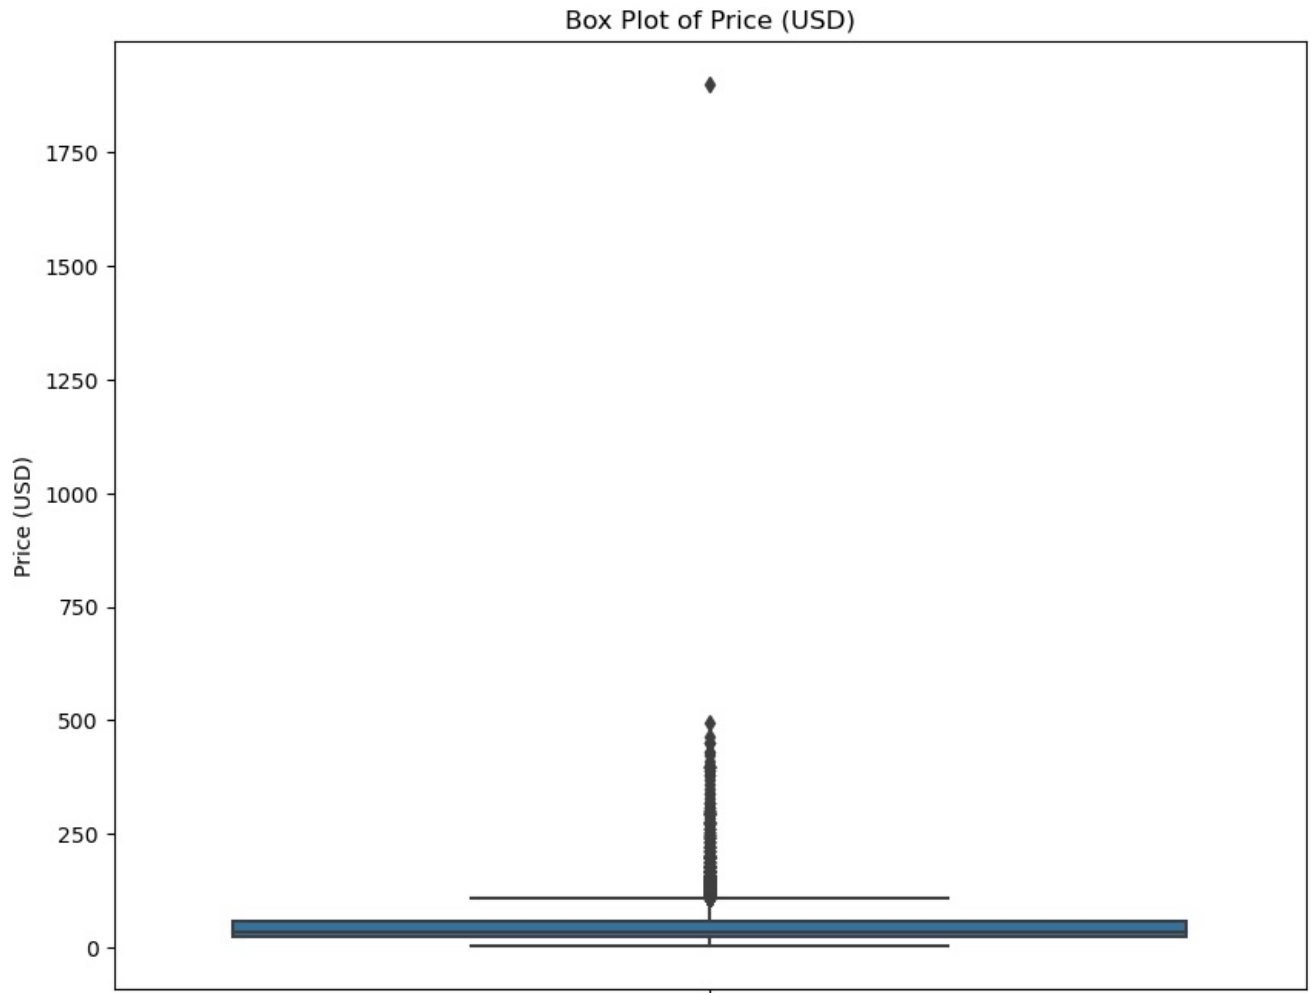
## INTERPRETATIONS:

From this, we can make note of the columns with a high percentage of missing data for when we assess which columns to remove from our dataset

```
In [45]: import seaborn as sns
import matplotlib.pyplot as plt

# Assuming 'data' is the DataFrame containing the data
# Replace 'data' with the actual DataFrame name

plt.figure(figsize=(10, 8))
sns.boxplot(data=data, y='price_usd')
plt.title('Box Plot of Price (USD)')
plt.ylabel('Price (USD)')
plt.show()
```



### Explanation:

By creating a box plot of the 'price\_usd' column, we can quickly observe the spread and central tendency of the prices in the dataset. It allows us to identify potential outliers, understand the range of prices, and assess the overall distribution of the prices in terms of quartiles.

## INTERPRETATION:

In summary, the box plot provides an effective and concise visual summary of the distribution of the 'price\_usd' column in the 'data' DataFrame, helping us better understand the pricing patterns and any potential outliers in the dataset.

```
In [46]: data.loc[data['price_usd']>1750]
```

```
Out[46]:
```

	product_id	product_name	brand_id	brand_name	loves_count	rating	reviews	size	variation_type	variation_value	...	online_only
6802	P502216	Shani Darden by Déesse PRO LED Light Mask	6314	Darden Skin Care	4154	3.75	4.0	NaN	NaN	NaN	...	1

1 rows × 27 columns

## INTERPRETATIONS:

The code `data.loc[data['price_usd']>1750]` retrieves rows from the data DataFrame where the value in the 'price\_usd' column is greater than 1750.

After doing some digging on Sephora's site, we can confirm that the price of the product above is legitimate. However, we will still exclude the outlier from the data to better gauge the price distribution later on.

Dropping the outlier

```
In [47]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "data"
# Replace "data" with the actual name of your DataFrame if different.

# Code to filter and drop the rows
data.drop(data[data['price_usd'] > 1750].index, inplace=True)

# The above code will drop the rows where 'price_usd' is greater than 1750.
```

```
In [48]: import pandas as pd

# Assuming your dataset is stored in a variable called "data"
# Replace "data" with the actual name of your DataFrame if different.

# Your code to read or import the dataset goes here...
# For example: data = pd.read_csv('your_dataset.csv')

# Drop the specified columns
columns_to_drop = [
    'loves_count',
    'tertiary_category',
    'highlights',
    'child_count',
    'child_max_price',
    'child_min_price',
    'sale_price_usd',
    'value_price_usd',
    'variation_value',
    'variation_desc'
]
data_edited = data.drop(columns=columns_to_drop)

# After dropping the columns, you can use "data_edited" instead of "data" for the updated dataset.
```

## INTERPRETATIONS:

We are dropping certain columns from the original dataset to create a new DataFrame called "data\_edited." The reason for doing this is to remove unnecessary or irrelevant columns that are not needed for the analysis or that contain a large number of missing values. By dropping these columns, we can create a more focused and compact dataset that only contains the relevant information, making it easier to work with and potentially improving the efficiency of any subsequent data analysis or modeling tasks.

```
In [49]: data_edited.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 8493 entries, 0 to 8493
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   product_id            8493 non-null   object
1   product_name          8493 non-null   object
2   brand_id              8493 non-null   int64
3   brand_name            8493 non-null   object
4   rating                8215 non-null   float64
5   reviews              8215 non-null   float64
6   size                  6863 non-null   object
7   variation_type        7050 non-null   object
8   ingredients            7549 non-null   object
9   price_usd             8493 non-null   float64
10  limited_edition       8493 non-null   int64
11  new                   8493 non-null   int64
12  online_only           8493 non-null   int64
13  out_of_stock          8493 non-null   int64
14  sephora_exclusive     8493 non-null   int64
15  primary_category      8493 non-null   object
16  secondary_category    8485 non-null   object
dtypes: float64(3), int64(6), object(8)
memory usage: 1.2+ MB
```

```
In [50]: data_edited.shape
```

```
Out[50]: (8493, 17)
```


We are creating a heatmap of null values in the "data\_edited" DataFrame to visualize and understand the presence of missing data within the dataset.

```
In [51]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Create a DataFrame to store information about null values
null_heatmap_data = data_edited.isnull()

# Create the heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(null_heatmap_data, cbar=False, cmap='viridis')
plt.title('Null Value Heatmap')
plt.show()
```



## INTERPRETATIONS:

By visualizing the null values using a heatmap, we can make informed decisions on how to handle missing data, such as imputing missing values, removing rows or columns with too many missing values, or applying specific data cleaning techniques.

```
In [52]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Calculate the mode for 'rating' and 'reviews'
rating_mode = data_edited['rating'].mode().iloc[0]
reviews_mode = data_edited['reviews'].mode().iloc[0]

# Fill the null values with the calculated modes
data_edited['rating'].fillna(rating_mode, inplace=True)
data_edited['reviews'].fillna(reviews_mode, inplace=True)
```

## INTERPRETATIONS:

Filling missing values with the mode is a common data imputation technique used when dealing with missing data in a dataset. By using the mode, we are essentially replacing the missing values with the most frequently occurring values in the respective columns.

```
In [53]: import pandas as pd
```



```

# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Get unique values in the 'variation_type' column
unique_variation_types = data_edited['variation_type'].unique()

# Print the unique values
print(unique_variation_types)

[nan 'Size + Concentration + Formulation' 'Scent' 'Size' 'Color'
 'Size + Concentration' 'Type' 'Formulation']

```

In [54]: `import pandas as pd`

```

# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Get unique values in the 'size' column
unique_sizes = data_edited['size'].unique()

# Print the unique values
print(unique_sizes)

[nan '3.4 oz/ 100 mL' '0.25 oz/ 7.5 mL' ... '0.25 oz/ 7.5 mL' '2.6 oz'
 '.11 oz / 3.2 mL']

```

In [55]: `data.isnull().sum()`

```

Out[55]: product_id          0
product_name        0
brand_id            0
brand_name          0
loves_count         0
rating              278
reviews             278
size                1630
variation_type      1443
variation_value     1597
variation_desc      7243
ingredients         944
price_usd           0
value_price_usd     8042
sale_price_usd      8223
limited_edition      0
new                 0
online_only         0
out_of_stock        0
sephora_exclusive   0
highlights          2207
primary_category    0
secondary_category   8
tertiary_category   990
child_count         0
child_max_price     5739
child_min_price     5739
dtype: int64

```

In [56]: `import pandas as pd`

```

# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Calculate the mode for 'secondary_category'
secondary_category_mode = data_edited['secondary_category'].mode().iloc[0]

# Fill null values with the calculated mode for 'secondary_category'
data_edited['secondary_category'].fillna(secondary_category_mode, inplace=True)

# Fill null values in 'size' and 'variation_type' with 'Not Available'
data_edited['size'].fillna('Not Available', inplace=True)
data_edited['variation_type'].fillna('Not Available', inplace=True)

```

In [57]: `import pandas as pd`

```

# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Get unique values in the 'size' column
unique_sizes = data_edited['size'].unique()

# Print the unique values
print(unique_sizes)

['Not Available' '3.4 oz/ 100 mL' '0.25 oz/ 7.5 mL' ... '0.25 oz/ 7.5 mL'
 '2.6 oz' '.11 oz / 3.2 mL']

```

In [58]: `import pandas as pd`

```
# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Get unique values in the 'variation_type' column
unique_variation_types = data_edited['variation_type'].unique()

# Print the unique values
print(unique_variation_types)

['Not Available' 'Size + Concentration + Formulation' 'Scent' 'Size'
 'Color' 'Size + Concentration' 'Type' 'Formulation']
```

In [59]: `import pandas as pd`

```
# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Fill null values in 'ingredients' column with 'Not available'
data_edited['ingredients'].fillna('Not available', inplace=True)
```

In [60]: `data_edited.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8493 entries, 0 to 8493
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   product_id            8493 non-null  object
1   product_name          8493 non-null  object
2   brand_id              8493 non-null  int64
3   brand_name            8493 non-null  object
4   rating                8493 non-null  float64
5   reviews               8493 non-null  float64
6   size                  8493 non-null  object
7   variation_type        8493 non-null  object
8   ingredients            8493 non-null  object
9   price_usd             8493 non-null  float64
10  limited_edition       8493 non-null  int64
11  new                   8493 non-null  int64
12  online_only           8493 non-null  int64
13  out_of_stock          8493 non-null  int64
14  sephora_exclusive     8493 non-null  int64
15  primary_category      8493 non-null  object
16  secondary_category    8493 non-null  object
dtypes: float64(3), int64(6), object(8)
memory usage: 1.2+ MB
```

In [61]: `# Assuming you have already defined the variable data_edited with the edited dataset`

```
numeric_cols = data_edited.select_dtypes(include=['number']).columns
print(numeric_cols)
print(f'{len(numeric_cols)} Numeric Columns in Products Dataset')
```

```
Index(['brand_id', 'rating', 'reviews', 'price_usd', 'limited_edition', 'new',
       'online_only', 'out_of_stock', 'sephora_exclusive'],
      dtype='object')
9 Numeric Columns in Products Dataset
```

In [62]: `# Assuming you have already defined the variable data_edited with the edited dataset`

```
numeric_cols = data_edited.select_dtypes(include=['number']).columns
print(numeric_cols)
print(f'{len(numeric_cols)} Numeric Columns in Products Dataset')
```

```
Index(['brand_id', 'rating', 'reviews', 'price_usd', 'limited_edition', 'new',
       'online_only', 'out_of_stock', 'sephora_exclusive'],
      dtype='object')
9 Numeric Columns in Products Dataset
```

In [63]: `# Assuming you have a DataFrame named 'data_edited'`

```
# Columns to convert from numeric (int64) to object (string)
columns_to_convert = ['limited_edition', 'new', 'online_only', 'out_of_stock', 'sephora_exclusive']

# Convert the selected columns to object (string) dtype
data_edited[columns_to_convert] = data_edited[columns_to_convert].astype(str)

# Verify the new dtypes of the selected columns
print(data_edited.dtypes)
```

```
product_id      object
product_name    object
brand_id        int64
brand_name      object
rating          float64
reviews         float64
size            object
variation_type  object
ingredients      object
price_usd       float64
limited_edition  object
new             object
online_only     object
out_of_stock    object
sephora_exclusive object
primary_category object
secondary_category object
dtype: object
```

```
In [64]: data_edited.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8493 entries, 0 to 8493
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   product_id            8493 non-null   object
1   product_name          8493 non-null   object
2   brand_id               8493 non-null   int64
3   brand_name            8493 non-null   object
4   rating                8493 non-null   float64
5   reviews               8493 non-null   float64
6   size                  8493 non-null   object
7   variation_type        8493 non-null   object
8   ingredients            8493 non-null   object
9   price_usd             8493 non-null   float64
10  limited_edition       8493 non-null   object
11  new                   8493 non-null   object
12  online_only           8493 non-null   object
13  out_of_stock          8493 non-null   object
14  sephora_exclusive     8493 non-null   object
15  primary_category      8493 non-null   object
16  secondary_category     8493 non-null   object
dtypes: float64(3), int64(1), object(13)
memory usage: 1.2+ MB
```

```
In [65]: data_edited.head(10)
```

Out[65]:

	product_id	product_name	brand_id	brand_name	rating	reviews	size	variation_type	ingredients	price_usd	limited_edition	new
0	P473671	Fragrance Discovery Set	6342	19-69	3.6364	11.0	Not Available	Not Available	['Capri Eau de Parfum:', 'Alcohol Denat. (SD A...	35.0	0	0
1	P473668	La Habana Eau de Parfum	6342	19-69	4.1538	13.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	0	0
2	P473662	Rainbow Bar Eau de Parfum	6342	19-69	4.2500	16.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	0	0
3	P473660	Kasbah Eau de Parfum	6342	19-69	4.4762	21.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	0	0
4	P473658	Purple Haze Eau de Parfum	6342	19-69	3.2308	13.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	0	0
5	P473661	Kasbah Eau de Parfum Travel Spray	6342	19-69	4.4762	21.0	0.25 oz/ 7.5 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	30.0	0	0
6	P473659	Purple Haze Eau de Parfum Travel Spray	6342	19-69	3.2308	13.0	0.25 oz/ 7.5 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	30.0	0	0
7	P473666	Invisible Post Eau de Parfum	6342	19-69	3.6250	8.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	0	0
8	P472300	Capri Eau de Parfum	6342	19-69	3.5714	7.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	0	0
9	P473667	Invisible Post Eau de Parfum Travel Spray	6342	19-69	3.6250	8.0	0.25 oz/ 7.5 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	30.0	0	0

In [66]:

data\_edited

Out[66]:

	product_id	product_name	brand_id	brand_name	rating	reviews	size	variation_type	ingredients	price_usd	limited_edit
0	P473671	Fragrance Discovery Set	6342	19-69	3.6364	11.0	Not Available	Not Available	['Capri Eau de Parfum:', 'Alcohol Denat. (SD A...	35.0	
1	P473668	La Habana Eau de Parfum	6342	19-69	4.1538	13.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	
2	P473662	Rainbow Bar Eau de Parfum	6342	19-69	4.2500	16.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	
3	P473660	Kasbah Eau de Parfum	6342	19-69	4.4762	21.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	
4	P473658	Purple Haze Eau de Parfum	6342	19-69	3.2308	13.0	3.4 oz/ 100 mL	Size + Concentration + Formulation	['Alcohol Denat. (SD Alcohol 39C), Parfum (Fra...	195.0	
...	...	...	...	...	...	...	...	...	...	...	
8489	P467659	Couture Clutch Eyeshadow Palette	1070	Yves Saint Laurent	4.4286	7.0	Not Available	Not Available	['Talc, Synthetic Fluorophlogopite, Triethylhex...	150.0	
8490	P500874	L'Homme Eau de Parfum	1070	Yves Saint Laurent	4.6367	556.0	2 oz / 60 mL	Size + Concentration + Formulation	['Alcohol, Aqua / Water / Eau, Parfum / Fragra...	106.0	
8491	P504428	Mon Paris Eau de Parfum Gift Set	1070	Yves Saint Laurent	5.0000	2.0	Not Available	Not Available	['Mon Paris Eau de Parfum:', 'Alcohol, Parfum/...	134.0	
8492	P504448	Y Eau de Parfum Gift Set	1070	Yves Saint Laurent	5.0000	2.0	Not Available	Not Available	['Alcohol, Parfum/Fragrance, Aqua/Water, Limon...	167.0	
8493	P505461	Candy Glaze Lip Gloss Stick Duo with Hyaluroni...	1070	Yves Saint Laurent	5.0000	2.0	.11 oz / 3.2 mL	Color	['Diisostearyl Malate, Bis- Behenyl/Isostearyl/...	50.0	

8493 rows × 17 columns

In [67]:

```
data_edited.to_csv("C:\\Users\\logeshwar\\Downloads\\Sephora Product dataset.csv",index=False)
```

## Data Cleaning on Reviews Dataset

Here we use the following code to see the amount of nulls by column within the dataset:

In [68]:

```
texta.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1301136 entries, 0 to 49976
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            1301136 non-null  int64
1   author_id                             1301136 non-null  object
2   rating                                1301136 non-null  int64
3   is_recommended                        1107162 non-null  float64
4   helpfulness                           631670 non-null  float64
5   total_feedback_count                  1301136 non-null  int64
6   total_neg_feedback_count              1301136 non-null  int64
7   total_pos_feedback_count              1301136 non-null  int64
8   submission_time                       1301136 non-null  object
9   review_text                           1299520 non-null  object
10  review_title                           930754 non-null  object
11  skin_tone                             1103798 non-null  object
12  eye_color                             1057734 non-null  object
13  skin_type                             1172830 non-null  object
14  hair_color                            1037824 non-null  object
15  product_id                            1301136 non-null  object
16  product_name                          1301136 non-null  object
17  brand_name                            1301136 non-null  object
18  price_usd                             1301136 non-null  float64
dtypes: float64(3), int64(5), object(11)
memory usage: 198.5+ MB
```

In [69]:

```
texta.shape
```

Out[69]:

(1301136, 19)

In [70]:

```
# heatmap to visualize missing data (reviews)


import seaborn as sns
```

```
import matplotlib.pyplot as plt

plt.figure(figsize=(10,8))

cols= texta.columns
colors=['black','purple']
sns.heatmap(texta[cols].isna(),cmap=sns.color_palette(colors))
```

Out[70]: <AxesSubplot:>



## INTERPRETATIONS:

From this, we can make note of the columns with a high percentage of missing data for when we assess which columns to remove from our dataset

```
In [71]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "texta"
# Replace "texta" with the actual name of your DataFrame if different.

# Get unique values in the 'helpfulness' column
unique_helpfulness = texta['helpfulness'].unique()

# Print the unique values
print(unique_helpfulness)
```

[1. nan 0.25 ... 0.182927 0.265625 0.90131599]

## Dropping the unnecessary columns

```
In [72]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "text_edited"
```

```
# Replace "text_edited" with the actual name of your DataFrame if different.

# Drop the 'helpfulness' column and the 'Unnamed: 0' column
texta = texta.drop(['helpfulness', 'Unnamed: 0'], axis=1)

# After dropping the columns, 'helpfulness' and 'Unnamed: 0', you can use "text_edited" for the updated dataset
```


```
In [73]: # Change the DataFrame name from "texta" to "text_edited"
text_edited = texta
```

```
In [74]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming your DataFrame is stored in a variable called "text_edited"
# Replace "text_edited" with the actual name of your DataFrame if different.

# Create a DataFrame to store information about null values
null_heatmap_data = text_edited.isnull()

# Create the heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(null_heatmap_data, cbar=False, cmap='viridis')
plt.title('Null Value Heatmap')
plt.show()
```



```
In [75]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "text edited"
# Replace "text_edited" with the actual name of your DataFrame if different.

# Calculate the modes for 'eye_color', 'skin_type', 'hair_color', and 'skin_tone'
eye_color_mode = text_edited['eye_color'].mode().iloc[0]
skin_type_mode = text_edited['skin_type'].mode().iloc[0]
hair_color_mode = text_edited['hair_color'].mode().iloc[0]
skin_tone_mode = text_edited['skin_tone'].mode().iloc[0]

# Fill null values with the calculated modes
text_edited['eye_color'].fillna(eye_color_mode, inplace=True)
text_edited['skin_type'].fillna(skin_type_mode, inplace=True)
text_edited['hair_color'].fillna(hair_color_mode, inplace=True)
```

```
text_edited['skin_tone'].fillna(skin_tone_mode, inplace=True)
```

```
In [76]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "text_edited"
# Replace "text_edited" with the actual name of your DataFrame if different.

# Fill null values in 'review_text' and 'review_title' with "Not Available"
text_edited['review_text'].fillna("Not Available", inplace=True)
text_edited['review_title'].fillna("Not Available", inplace=True)
```

```
In [77]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "text_edited"
# Replace "text_edited" with the actual name of your DataFrame if different.

# Get unique values in the 'is_recommended' column
unique_is_recommended = text_edited['is_recommended'].unique()

# Print the unique values
print(unique_is_recommended)

[ 1.  0. nan]
```

```
In [78]: text_edited['is_recommended'].value_counts()
```

```
Out[78]: 1.0    929476
0.0    177686
Name: is_recommended, dtype: int64
```

```
In [79]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "text_edited"
# Replace "text_edited" with the actual name of your DataFrame if different.

# Calculate the mode for 'is_recommended'
is_recommended_mode = text_edited['is_recommended'].mode().iloc[0]

# Fill null values in 'is_recommended' with the calculated mode
text_edited['is_recommended'].fillna(is_recommended_mode, inplace=True)
```

```
In [80]: import pandas as pd

# Assuming your DataFrame is stored in a variable called "text_edited"
# Replace "text_edited" with the actual name of your DataFrame if different.

# Get unique values in the 'is_recommended' column
unique_is_recommended = text_edited['is_recommended'].unique()

# Print the unique values
print(unique_is_recommended)

[1. 0.]
```

```
In [81]: # Assuming you have already defined the variable text_edited with the edited dataset
numeric_cols_reviews = text_edited.select_dtypes(include=['number']).columns
print(numeric_cols_reviews)
print(f'{len(numeric_cols_reviews)} Numeric Columns in Reviews Dataset')

Index(['rating', 'is_recommended', 'total_feedback_count',
       'total_neg_feedback_count', 'total_pos_feedback_count', 'price_usd'],
      dtype='object')
6 Numeric Columns in Reviews Dataset
```

```
In [82]: import pandas as pd

# Assuming you have a DataFrame named 'text_edited'

# Convert 'is_recommended' column from float64 to object (string)
text_edited['is_recommended'] = text_edited['is_recommended'].astype(str)

# Convert 'submission_time' column to datetime64[ns]
text_edited['submission_time'] = pd.to_datetime(text_edited['submission_time'])

# Verify the new dtypes of the columns
print(text_edited.dtypes)
```



author\_id object  
rating int64  
is\_recommended object  
total\_feedback\_count int64  
total\_neg\_feedback\_count int64  
total\_pos\_feedback\_count int64  
submission\_time datetime64[ns]  
review\_text object  
review\_title object  
skin\_tone object  
eye\_color object  
skin\_type object  
hair\_color object  
product\_id object  
product\_name object  
brand\_name object  
price\_usd float64  
dtype: object

```
In [83]: text_edited.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 1301136 entries, 0 to 49976  
Data columns (total 17 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   author_id                            1301136 non-null object  
1   rating                               1301136 non-null int64  
2   is_recommended                       1301136 non-null object  
3   total_feedback_count                 1301136 non-null int64  
4   total_neg_feedback_count             1301136 non-null int64  
5   total_pos_feedback_count             1301136 non-null int64  
6   submission_time                      1301136 non-null datetime64[ns]  
7   review_text                          1301136 non-null object  
8   review_title                         1301136 non-null object  
9   skin_tone                           1301136 non-null object  
10  eye_color                            1301136 non-null object  
11  skin_type                            1301136 non-null object  
12  hair_color                           1301136 non-null object  
13  product_id                           1301136 non-null object  
14  product_name                         1301136 non-null object  
15  brand_name                           1301136 non-null object  
16  price_usd                            1301136 non-null float64  
dtypes: datetime64[ns](1), float64(1), int64(4), object(11)  
memory usage: 178.7+ MB
```

```
In [84]: text_edited.head()
```

Out[84]:

	author_id	rating	is_recommended	total_feedback_count	total_neg_feedback_count	total_pos_feedback_count	submission_time	review
--	-----------	--------	----------------	----------------------	--------------------------	--------------------------	-----------------	--------

0	1741593524	5	1.0	2	0	2	2023-02-01	I us wit Nuc "Ç Clean
1	31423088263	1	0.0	0	0	0	2023-03-21	I bc th mask readin re
2	5061282401	5	1.0	0	0	0	2023-03-21	My re title s all! I g excit
3	6083038851	5	1.0	0	0	0	2023-03-20	I've al love formu a ti
4	47056667835	5	1.0	0	0	0	2023-03-20	If you dry cra lips, thi mu:

```
In [85]: text_edited
```

Out[85]:

	author_id	rating	is_recommended	total_feedback_count	total_neg_feedback_count	total_pos_feedback_count	submission_time	re
0	1741593524	5	1.0	2	0	2	2023-02-01	C
1	31423088263	1	0.0	0	0	0	2023-03-21	i n
2	5061282401	5	1.0	0	0	0	2023-03-21	t e
3	6083038851	5	1.0	0	0	0	2023-03-20	f
4	47056667835	5	1.0	0	0	0	2023-03-20	It di lip
...	...	...	...	...	...	...	...	...
49972	2276253200	5	1.0	0	0	0	2023-03-13	
49973	28013163278	5	1.0	0	0	0	2023-03-13	a n
49974	1539813076	5	1.0	0	0	0	2023-03-13	it c
49975	5595682861	5	1.0	0	0	0	2023-03-13	I a
49976	27666075558	5	1.0	0	0	0	2023-03-13	r

1301136 rows × 17 columns

# EXPLORATORY DATA ANALYSIS

## EDA on Product Dataset

In [86]:

```
# Display basic information about the dataset
print(data_edited.info())
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 8493 entries, 0 to 8493
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   product_id            8493 non-null   object
1   product_name          8493 non-null   object
2   brand_id              8493 non-null   int64
3   brand_name            8493 non-null   object
4   rating                8493 non-null   float64
5   reviews              8493 non-null   float64
6   size                  8493 non-null   object
7   variation_type        8493 non-null   object
8   ingredients            8493 non-null   object
9   price_usd             8493 non-null   float64
10  limited_edition       8493 non-null   object
11  new                    8493 non-null   object
12  online_only           8493 non-null   object
13  out_of_stock          8493 non-null   object
14  sephora_exclusive     8493 non-null   object
15  primary_category      8493 non-null   object
16  secondary_category    8493 non-null   object
dtypes: float64(3), int64(1), object(13)
memory usage: 1.2+ MB
None

```

```

In [87]: # Summary statistics of numeric columns
print(data_edited.describe())

```

	brand_id	rating	reviews	price_usd
count	8493.000000	8493.000000	8493.000000	8493.000000
mean	5422.335570	4.220931	433.981161	51.437963
std	1709.669236	0.527999	1086.759660	49.783262
min	1063.000000	1.000000	1.000000	3.000000
25%	5333.000000	4.000000	22.000000	25.000000
50%	6156.000000	4.308800	112.000000	35.000000
75%	6328.000000	4.558400	402.000000	58.000000
max	8020.000000	5.000000	21281.000000	495.000000

```

In [88]: # Distribution of numerical features

```




```

# Histogram of 'rating' column
plt.hist(data_edited['rating'], bins=20, edgecolor='black')
plt.xlabel('Rating')
plt.ylabel('Frequency')
plt.title('Distribution of Ratings')
plt.show()

# Histogram of 'reviews' column
plt.hist(data_edited['reviews'], bins=20, edgecolor='black')
plt.xlabel('Number of Reviews')
plt.ylabel('Frequency')
plt.title('Distribution of Reviews')
plt.show()

# Histogram of 'price_usd' column
plt.hist(data_edited['price_usd'], bins=20, edgecolor='black')
plt.xlabel('Price (USD)')
plt.ylabel('Frequency')
plt.title('Distribution of Prices')
plt.show()

```



```
In [89]: import pandas as pd
import matplotlib.pyplot as plt
```

```
# Assuming your DataFrame is stored in a variable called "data_edited"
# Replace "data_edited" with the actual name of your DataFrame if different.

# Proportion of limited edition, new, online-only, out of stock, and Sephora exclusive products
limited_edition_counts = data_edited['limited_edition'].value_counts(normalize=True) * 100
new_counts = data_edited['new'].value_counts(normalize=True) * 100
online_only_counts = data_edited['online_only'].value_counts(normalize=True) * 100
out_of_stock_counts = data_edited['out_of_stock'].value_counts(normalize=True) * 100
sephora_exclusive_counts = data_edited['sephora_exclusive'].value_counts(normalize=True) * 100

# Plotting the proportions as bar plots
fig, axs = plt.subplots(2, 3, figsize=(15, 10))

# Bar plot for limited edition
axs[0, 0].bar(limited_edition_counts.index, limited_edition_counts.values)
axs[0, 0].set_title("Proportion of Limited Edition Products")
axs[0, 0].set_xlabel("Limited Edition")
axs[0, 0].set_ylabel("Proportion (%)")

# Bar plot for new products
axs[0, 1].bar(new_counts.index, new_counts.values)
axs[0, 1].set_title("Proportion of New Products")
axs[0, 1].set_xlabel("New Product")
axs[0, 1].set_ylabel("Proportion (%)")

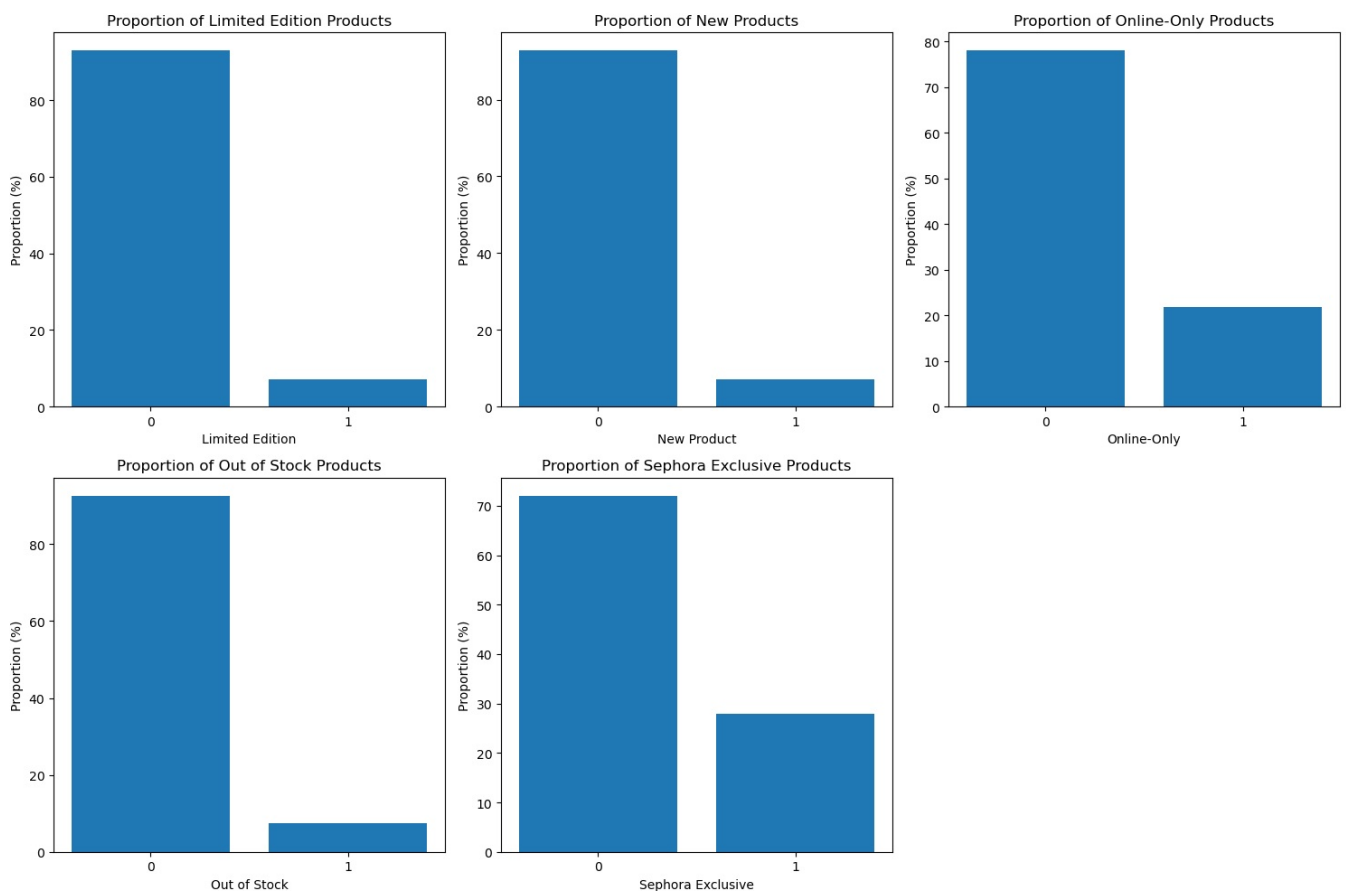
# Bar plot for online-only products
axs[0, 2].bar(online_only_counts.index, online_only_counts.values)
axs[0, 2].set_title("Proportion of Online-Only Products")
axs[0, 2].set_xlabel("Online-Only")
axs[0, 2].set_ylabel("Proportion (%)")

# Bar plot for out of stock products
axs[1, 0].bar(out_of_stock_counts.index, out_of_stock_counts.values)
axs[1, 0].set_title("Proportion of Out of Stock Products")
axs[1, 0].set_xlabel("Out of Stock")
axs[1, 0].set_ylabel("Proportion (%)")

# Bar plot for Sephora exclusive products
axs[1, 1].bar(sephora_exclusive_counts.index, sephora_exclusive_counts.values)
axs[1, 1].set_title("Proportion of Sephora Exclusive Products")
axs[1, 1].set_xlabel("Sephora Exclusive")
axs[1, 1].set_ylabel("Proportion (%)")

# Removing the empty subplot
fig.delaxes(axs[1, 2])

plt.tight_layout()
plt.show()
```



INTERPRETATION:

Most products aren't "Limited Edition" or "New".

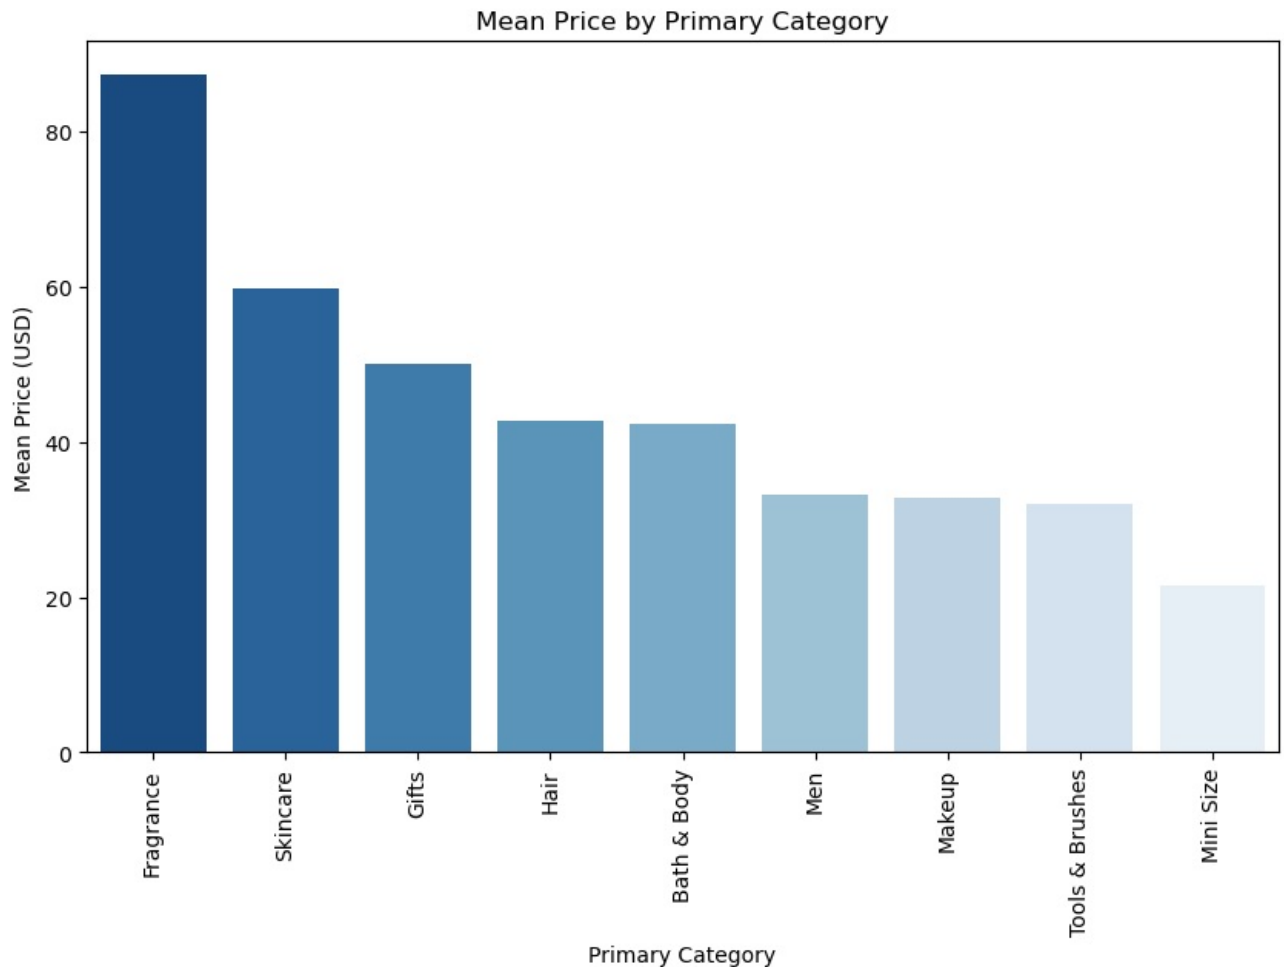
About 1 in 5 products are "Out of Stock".

Roughly 1 in 3 products are "Sephora Exclusive".

A quarter of the products are sold "Online-Only".

In short, most products are regular items, with a few being exclusive to Sephora or online-only.

```
In [90]: mean_price_by_category = data_edited.groupby('primary_category')['price_usd'].mean().sort_values(ascending=False)
plt.figure(figsize=(10, 6))
sns.barplot(x=mean_price_by_category.index, y=mean_price_by_category.values, palette='Blues_r')
plt.xlabel('Primary Category')
plt.ylabel('Mean Price (USD)')
plt.title('Mean Price by Primary Category')
plt.xticks(rotation=90)
plt.show()
```



## INTERPRETATION:

The graph shows the average price of products for different categories at Sephora. "Fragrance" has the highest average price, while "Mini Size" items have the lowest average price.

```
In [91]: import pandas as pd
import matplotlib.pyplot as plt

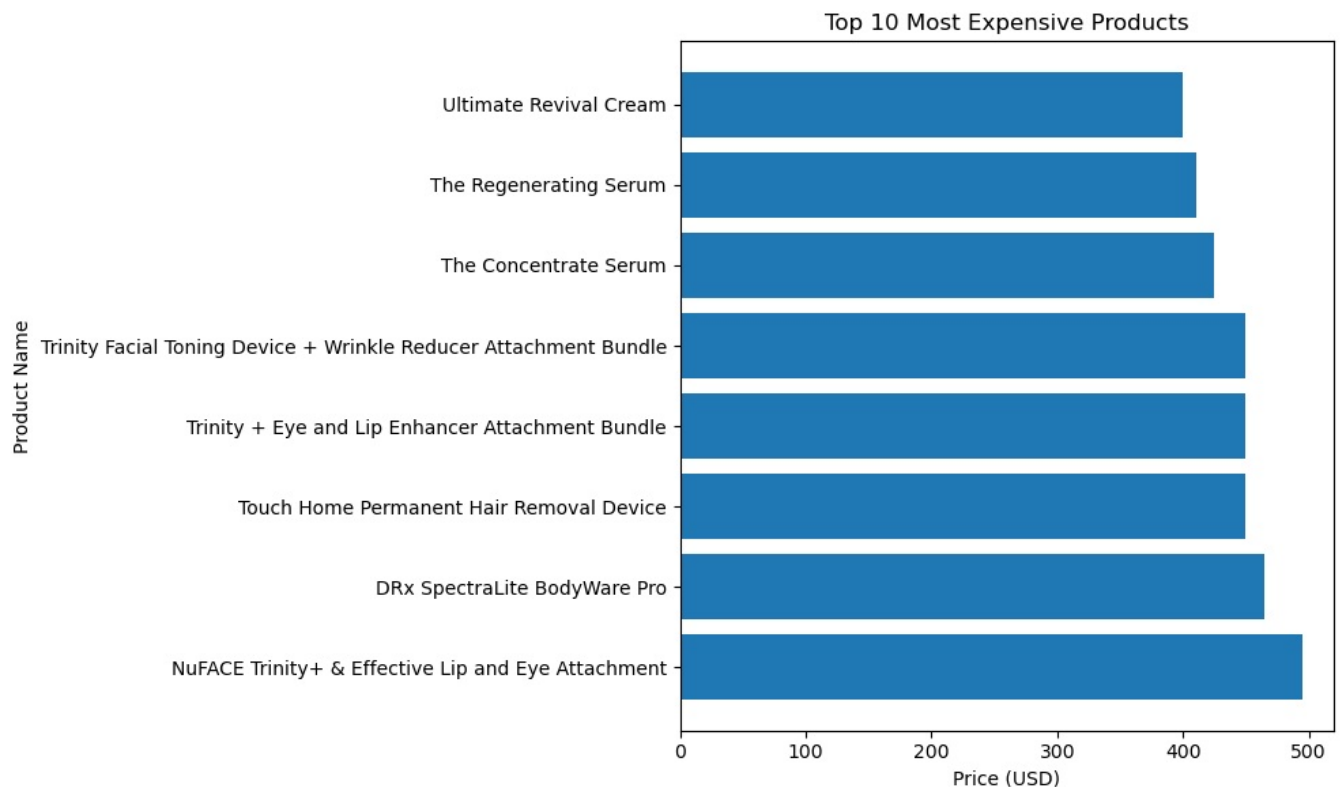
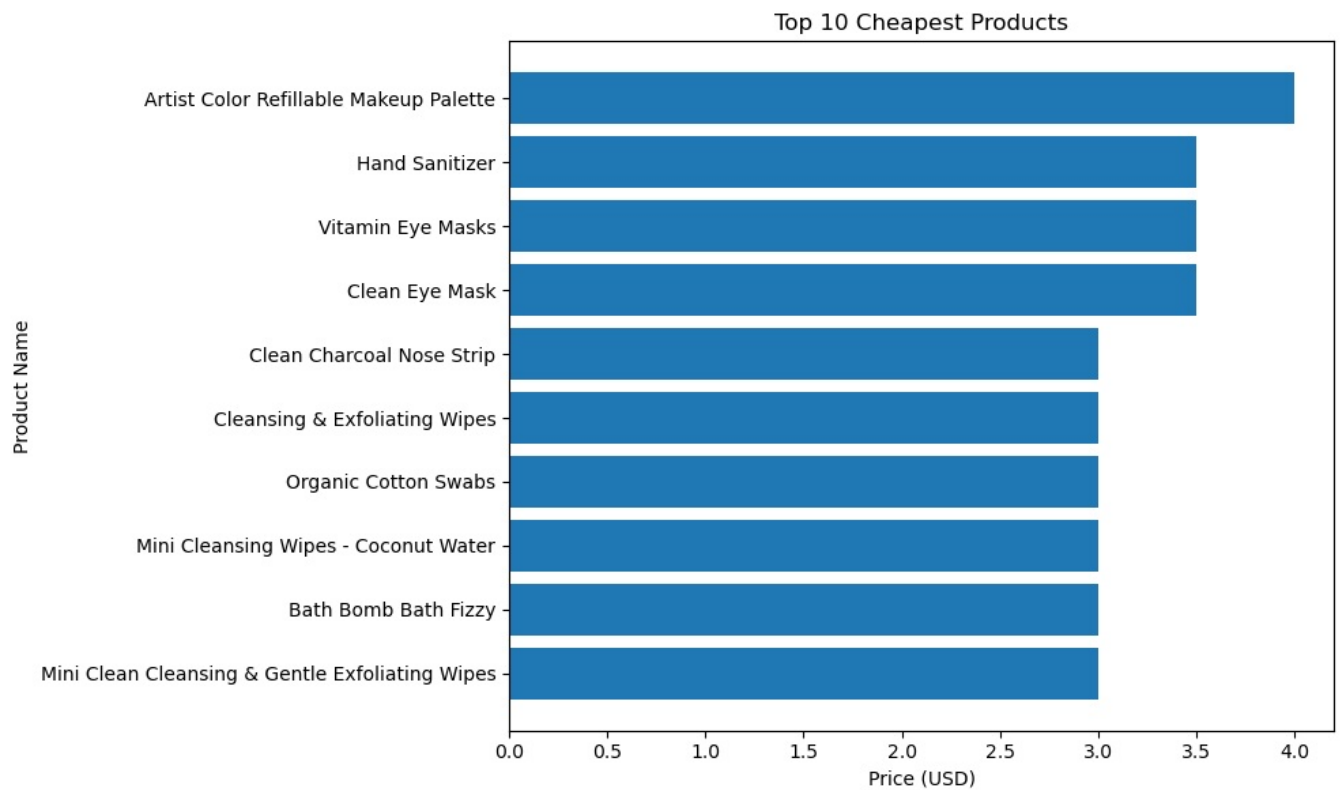
# Assuming you have a DataFrame named 'data_edited'

# Sort the DataFrame by 'price_usd' in ascending order to get cheapest products first
cheapest_products = data_edited.sort_values(by='price_usd').head(10)

# Sort the DataFrame by 'price_usd' in descending order to get most expensive products first
expensive_products = data_edited.sort_values(by='price_usd', ascending=False).head(10)

# Create a bar chart for the top 10 cheapest products
plt.figure(figsize=(10, 6))
plt.barh(cheapest_products['product_name'], cheapest_products['price_usd'])
plt.title('Top 10 Cheapest Products')
plt.xlabel('Price (USD)')
plt.ylabel('Product Name')
plt.tight_layout()
plt.show()
```

```
# Create a bar chart for the top 10 most expensive products
plt.figure(figsize=(10, 6))
plt.barh(expensive_products['product_name'], expensive_products['price_usd'])
plt.title('Top 10 Most Expensive Products')
plt.xlabel('Price (USD)')
plt.ylabel('Product Name')
plt.tight_layout()
plt.show()
```



## INTERPRETATION:

The graphs display the top 10 cheapest and most expensive products at Sephora.

**Cheapest Products:** Items like "Hand Sanitizer" and "Clean Eye Mask" are among the least expensive, with prices under \$4. **Most Expensive Products:** Products like "Ultimate Revival Cream" and "The Regenerating Serum" are among the priciest, costing several hundred dollars each. In essence, Sephora offers a wide price range, from very affordable items to premium-priced products.

## EDA on Reviews Dataset

```


<class 'pandas.core.frame.DataFrame'>
Int64Index: 1301136 entries, 0 to 49976
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype  
---  -
0   author_id                             1301136 non-null object  
1   rating                                1301136 non-null int64   
2   is_recommended                        1301136 non-null object  
3   total_feedback_count                 1301136 non-null int64   
4   total_neg_feedback_count             1301136 non-null int64   
5   total_pos_feedback_count             1301136 non-null int64   
6   submission_time                      1301136 non-null datetime64[ns]
7   review_text                          1301136 non-null object  
8   review_title                         1301136 non-null object  
9   skin_tone                            1301136 non-null object  
10  eye_color                            1301136 non-null object  
11  skin_type                            1301136 non-null object  
12  hair_color                           1301136 non-null object  
13  product_id                           1301136 non-null object  
14  product_name                         1301136 non-null object  
15  brand_name                           1301136 non-null object  
16  price_usd                            1301136 non-null float64  
dtypes: datetime64[ns](1), float64(1), int64(4), object(11)
memory usage: 178.7+ MB

```

```

In [93]: # Distribution plot of 'rating' column
plt.figure(figsize=(8, 6))
sns.histplot(text_edited['rating'], kde=True)
plt.xlabel('Rating')
plt.ylabel('Frequency')
plt.title('Distribution of Ratings')
plt.show()

```



## INTERPRETATION:

The graph showcases the distribution of product ratings:

A very small number of products received ratings around 1.0 to 3.0. The majority of products received the highest rating of 5.0. In essence, most products at Sephora have been highly rated by customers.

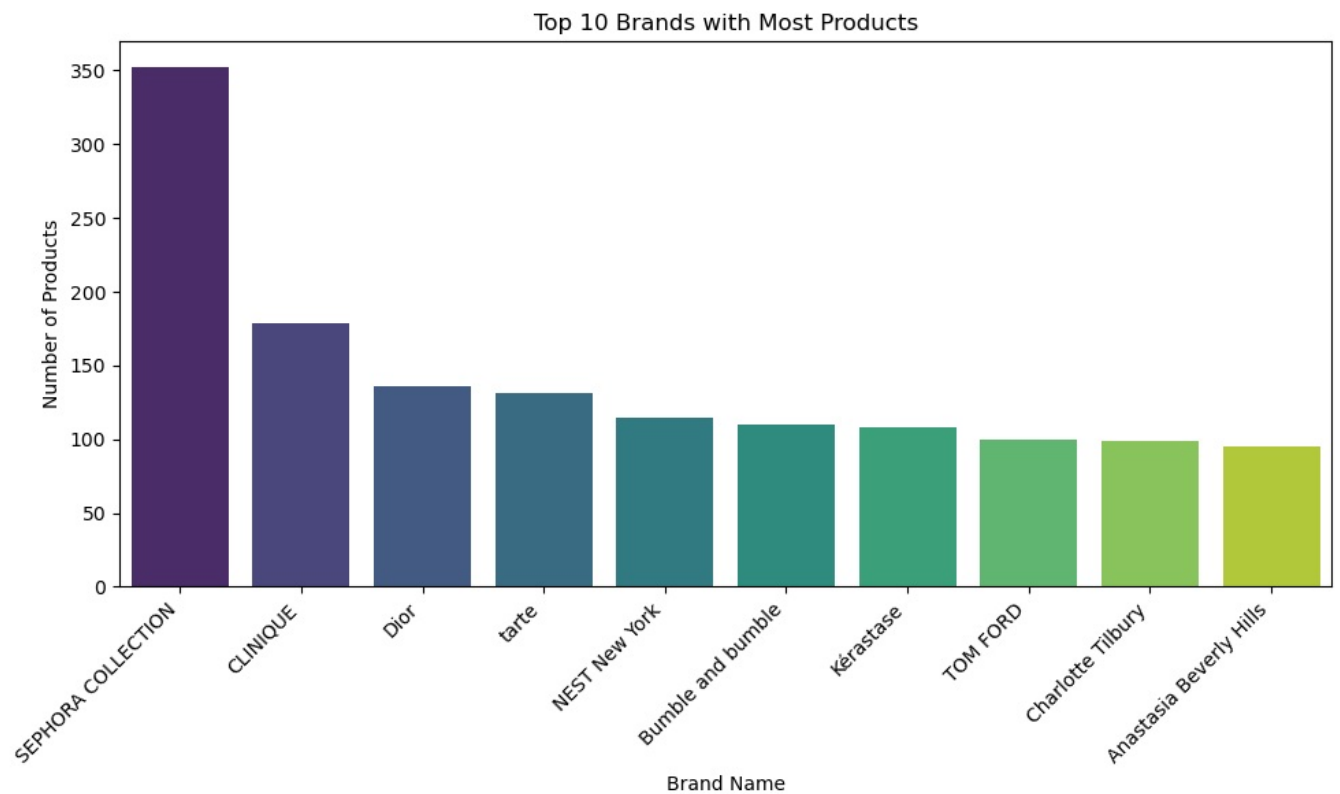
```

In [94]: top_brands = data['brand_name'].value_counts().nlargest(10)

plt.figure(figsize=(10, 6))
sns.barplot(x=top_brands.index, y=top_brands.values, palette='viridis')
plt.xticks(rotation=45, ha='right')
plt.xlabel('Brand Name')
plt.ylabel('Number of Products')
plt.title('Top 10 Brands with Most Products')
plt.tight_layout()
plt.show()

```





### INTERPRETATION:

The chart showcases the top 10 brands with the most products on Sephora.

SEPHORA COLLECTION has the highest number, over 300 products. CLINIQUE comes second with around 250 products. Brands like Dior, tarte, and NEST New York have between 100 to 200 products. The remaining brands, including TOM FORD and Charlotte Tilbury, offer 50 to 100 products each.

In essence, SEPHORA COLLECTION dominates in product variety, but several other brands also have a strong presence.

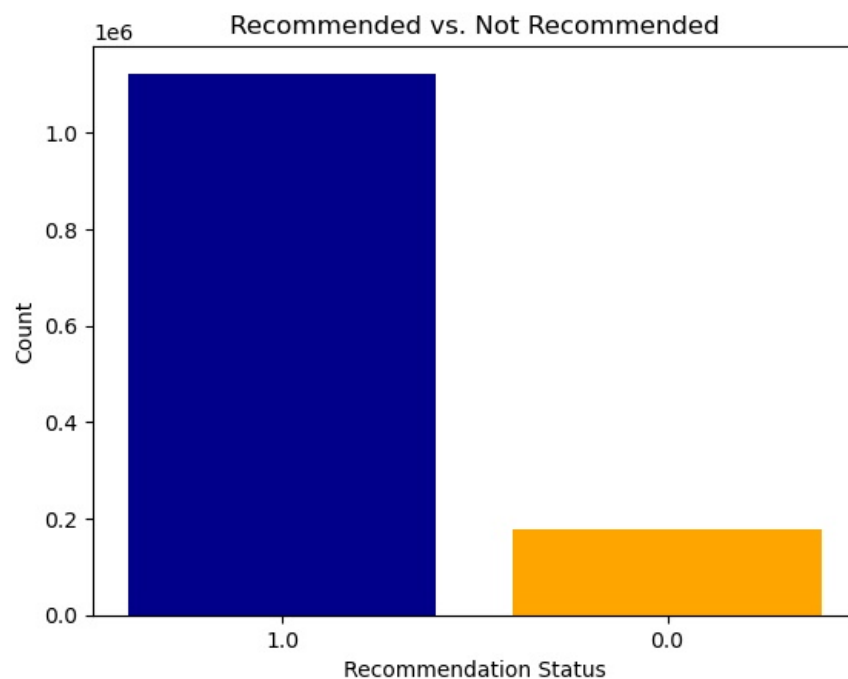
```
In [95]: import matplotlib.pyplot as plt
import pandas as pd

# Assuming your dataset is stored in a DataFrame called 'text_edited'
# If your dataset is in a CSV file, you can read it using pd.read_csv('filename.csv')

# Step 1: Count occurrences of each value in 'is_recommended' column
recommended_counts = text_edited['is_recommended'].value_counts()

# Step 2: Create the bar chart
plt.bar(recommended_counts.index, recommended_counts.values, color=['darkblue', 'orange'])

plt.xlabel('Recommendation Status')
plt.ylabel('Count')
plt.title('Recommended vs. Not Recommended')
plt.show()
```



The visualization titled "Recommended vs. Not Recommended" provides a clear comparison of the count of products that have been recommended versus those that haven't.


### INTERPRETATION:

A vast majority of products have been recommended, as indicated by the tall blue bar. This suggests that most users or reviewers had a positive experience with the products and found them satisfactory.

A significantly smaller number of products were not recommended, as shown by the shorter orange bar. This indicates that a lesser number of users had reservations or were not entirely satisfied with these products.

In summary, the visualization underscores a predominantly positive reception for products, with a small fraction not meeting the users' expectations or requirements.

```
In [96]: # Count of products by 'product_name'
plt.figure(figsize=(12, 6))
sns.barplot(y=text_edited['product_name'].value_counts().nlargest(10).index, x=text_edited['product_name'].value_counts().nlargest(10).values)
plt.ylabel('Product Name')
plt.xlabel('Count')
plt.title('Top 10 Most Reviewed Products')
plt.show()
```

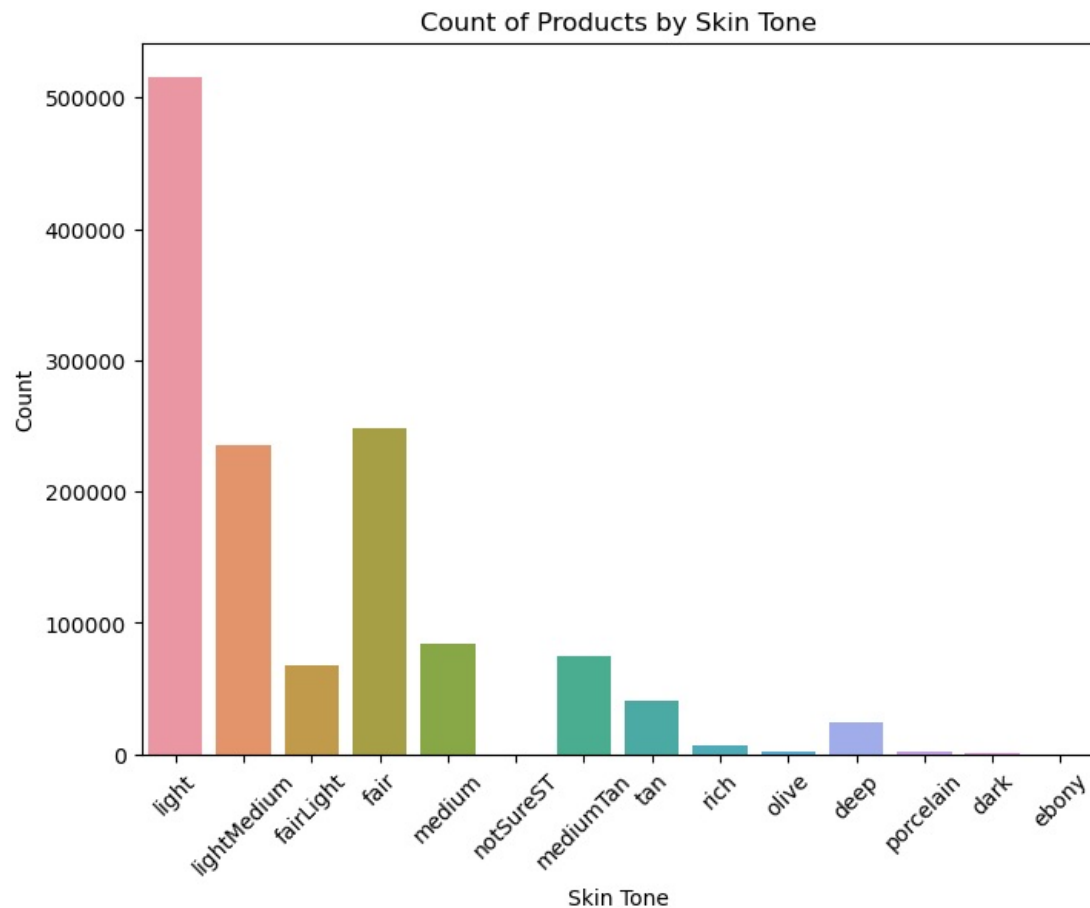


### INTERPRETATION:

The graph displays the top 10 products with the highest number of reviews. The "Lip Sleeping Mask Intense Hydration with Vitamin C" has the most reviews, followed by other products in decreasing order.

```
In [97]: # Count of Products by Skin Tone
plt.figure(figsize=(8, 6))
sns.countplot(x='skin_tone', data=text_edited)
plt.xlabel('Skin Tone')
plt.ylabel('Count')
plt.title('Count of Products by Skin Tone')
```

```
plt.xticks(rotation=45)
plt.show()
```



## INTERPRETATION:

The graph illustrates the count of products based on different skin tones. "Light" skin tone products have the highest count, followed by "light/medium" and "medium" skin tones. Other skin tones have comparatively fewer products.

## STATISTICAL TESTS:

### Test for Collinearity:

```
In [118.. import warnings
warnings.filterwarnings('ignore')
```

```
In [117.. import pandas as pd

# Load the dataset
data = pd.read_csv(r"C:\Users\logeshwar\Downloads\Sephora Reviews Dataset.csv")

# Calculate Pearson correlation coefficients for numerical variables
numerical_vars = ['rating', 'total_feedback_count', 'total_neg_feedback_count', 'total_pos_feedback_count', 'pr
correlation_matrix = data[numerical_vars].corr()
print("Correlation Matrix:")
print(correlation_matrix)
```

Correlation Matrix:

```
          rating  total_feedback_count  \
rating      1.000000      -0.080300
total_feedback_count -0.080300      1.000000
total_neg_feedback_count -0.182179      0.674619
total_pos_feedback_count -0.049147      0.984976
price_usd      -0.002616      0.008143

          total_neg_feedback_count  total_pos_feedback_count  \
rating      -0.182179      -0.049147
total_feedback_count      0.674619      0.984976
total_neg_feedback_count      1.000000      0.537009
total_pos_feedback_count      0.537009      1.000000
price_usd      0.007682      0.007508

          price_usd
rating      -0.002616
total_feedback_count      0.008143
total_neg_feedback_count      0.007682
total_pos_feedback_count      0.007508
price_usd      1.000000
```

## INTERPRETATION:

Products with higher ratings tend to have fewer negative feedbacks.

Products with more total feedback also have more positive and negative feedback.

Price doesn't show a clear relationship with ratings or feedback counts.

Basically, well-rated products get fewer negative comments, and popular products (with lots of feedback) have both more likes and dislikes. Product prices don't seem to affect ratings or the number of comments much.

## Chi-Square Test:

```
In [100] from scipy.stats import chi2_contingency

# Create a contingency table for two categorical variables
contingency_table = pd.crosstab(data['is_recommended'], data['skin_type'])

# Perform the chi-square test
chi2_stat, p_value, dof, expected = chi2_contingency(contingency_table)
print("Chi-Square Test - Chi2-statistic:", chi2_stat)
print("Chi-Square Test - P-value:", p_value)
```

```
Chi-Square Test - Chi2-statistic: 1923.860016218833
Chi-Square Test - P-value: 0.0
```

## INTERPRETATION:

The Chi-Square test was performed to check if there's a relationship between the recommendation status ('is\_recommended') of a product and the skin type of users ('skin\_type').

With a Chi2-statistic value of 1923.86 and a p-value of 0.0, the results are statistically significant. This suggests that there is a strong association between whether a product is recommended and the skin type of the users. In other words, the likelihood of a product being recommended may vary depending on the skin type of the user.

## Chi-square Test for Independence:

```
In [102] from scipy.stats import chi2_contingency

# List of categorical predictors including 'rating'
categorical_predictors = ['rating', 'skin_tone', 'eye_color', 'skin_type', 'hair_color']

# Running chi-square test for each predictor against 'is_recommended'
for predictor in categorical_predictors:
    crosstab = pd.crosstab(data[predictor], data['is_recommended'])
    chi2, p, _, _ = chi2_contingency(crosstab)
    print(f"P-value for {predictor} vs. is_recommended:", p)
```

```
P-value for rating vs. is_recommended: 0.0
P-value for skin_tone vs. is_recommended: 0.0
P-value for eye_color vs. is_recommended: 0.0
P-value for skin_type vs. is_recommended: 0.0
P-value for hair_color vs. is_recommended: 0.0
```

## INTERPRETATION:

All the p-values are 0.0, which means:

Rating, skin\_tone, eye\_color, skin\_type, and hair\_color each have a relationship with whether a product is recommended or not.

It's not just by random chance; there's a real association.

For example, products with higher ratings are more likely to be recommended. Similarly, people with different skin tones, eye colors, etc., might have varying preferences.


In short, all these factors play a role in whether a product gets recommended.

## Descriptive Statistics and Visualization:

```
In [103.. import matplotlib.pyplot as plt

data['price_usd'].hist()
plt.show()

print(data['is_recommended'].value_counts())
```



```
1.0    799381
0.0    124076
Name: is_recommended, dtype: int64
```

### INTERPRETATION:

The histogram shows how product prices are spread out.

Most products, around 86.5%, are recommended by customers, while the rest are not.

In short, the majority of products get a thumbs-up from customers, and the histogram shows their pricing distribution.

## T-Test:

```
In [104.. print(data['is_recommended'].unique())
print(data['is_recommended'].dtype)

[1. 0.]
float64
```

```
In [105.. from scipy.stats import ttest_ind

# Split the data into two groups: recommended and not recommended
recommended_prices = data[data['is_recommended'] == 1]['price_usd']
not_recommended_prices = data[data['is_recommended'] == 0]['price_usd']

# Perform the t-test
t_stat, p_value = ttest_ind(recommended_prices, not_recommended_prices)

print("T-statistic:", t_stat)
print("P-value:", p_value)
```

```
T-statistic: -2.9337376064389415
P-value: 0.0033491543629176807
```

### INTERPRETATION:

The T-statistic value of -2.9337 suggests that there's a difference between the two groups, with the average price of the not-recommended products being higher than the recommended ones (since the value is negative).

The P-value is 0.0033, which is less than the typical significance level of 0.05. This means that the difference between the average prices


of recommended and not-recommended products is statistically significant. In other words, the observed difference in the averages is likely not due to random chance.

In simple terms: The average price of products that aren't recommended is statistically higher than those that are recommended. make it simple

## Relationship between product prices (in USD) and their ratings.

```
In [107... import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 6))
sns.scatterplot(x='rating', y='price_usd', data=data, alpha=0.5)
plt.xlabel('Rating')
plt.ylabel('Price (USD)')
plt.title('Price vs. Rating')
plt.show()
```



The graph shows the relationship between product prices (in USD) and their ratings.

### INTERPRETATIONS:

Most products, regardless of their price, have a high rating (around 4 to 5).

There are some products with low ratings (1 to 2), but they are fewer in number.

Products with a wide range of prices can have similar high ratings.


Overall, there's no clear trend that higher-priced products have higher or lower ratings.

## Business Problem 1

Is there a significant difference in average product prices between products that are recommended and those that are not?

```
In [108... import matplotlib.pyplot as plt
import seaborn as sns

# Create a boxplot to visualize price differences
sns.boxplot(x='is_recommended', y='price_usd', data=data)
plt.title("Price Distribution by Recommendation Status")
plt.show()
```



### INTERPRETATION:

The boxplot displays the distribution of product prices based on their recommendation status. For products that are recommended (represented by 1.0), the price range appears broader, with both a higher median price and more outliers in the higher price range compared to products that are not recommended (represented by 0.0). The non-recommended products have a slightly lower median price, with fewer high-priced outliers.

## Solution to the Business Problem:

### Value for Money:

The wider range and slightly higher median price for recommended products suggest that customers are willing to pay a premium for products they perceive as valuable or of high quality. The brand should focus on emphasizing the value proposition and unique selling points of higher-priced products to maintain and even boost their recommendation status.

### Price Sensitivity:

The non-recommended products' lower price range hints at possible issues other than just the price – perhaps quality, packaging, or efficacy. It might be beneficial to revisit these products and analyze negative feedback to understand the core concerns.

### Review and Revise:

The presence of high-priced outliers in the non-recommended bracket suggests that there might be premium products that aren't meeting customer expectations. A deeper dive into the specific feedback for these products could offer insights into why they aren't resonating well with customers despite their high price point.

### Customer Engagement:

Engage with customers to understand what improvements can be made to non-recommended products, especially those in the higher price range. Feedback can be invaluable in making necessary changes to the product or its positioning.

### Pricing Strategy:

Consider periodic reviews of the pricing strategy, taking into account customer feedback and market trends. Offering promotions, discounts, or bundling options might help in improving the recommendation status of certain products.

In conclusion, while price plays a role in a product's recommendation status, it's crucial to understand that customers value quality, efficacy, and overall satisfaction. Ensuring that products, whether low or high priced, meet or exceed customer expectations will positively impact their recommendation status.

## Business Problem 2

"To launch a new marketing campaign, we want to identify which customer segments (based on features like skin tone, hair color, eye color) might be underrepresented or overly dissatisfied in our reviews. Identifying these groups will allow us to tailor our campaigns to these segments, ensuring they feel represented and addressing their specific concerns."

```
In [109.. import matplotlib.pyplot as plt
import seaborn as sns

# Distribution of reviews based on skin tone, hair color, and eye color
fig, axes = plt.subplots(nrows=3, ncols=1, figsize=(10, 15))

sns.countplot(x='skin_tone', data=data, ax=axes[0], palette='pastel')
axes[0].set_title('Distribution of Reviews by Skin Tone')
axes[0].set_ylabel('Number of Reviews')

sns.countplot(x='hair_color', data=data, ax=axes[1], palette='pastel')
axes[1].set_title('Distribution of Reviews by Hair Color')
axes[1].set_ylabel('Number of Reviews')

sns.countplot(x='eye_color', data=data, ax=axes[2], palette='pastel')
axes[2].set_title('Distribution of Reviews by Eye Color')
axes[2].set_ylabel('Number of Reviews')

plt.tight_layout()
plt.show()

# Average satisfaction (rating) for each segment
fig, axes = plt.subplots(nrows=3, ncols=1, figsize=(10, 15))

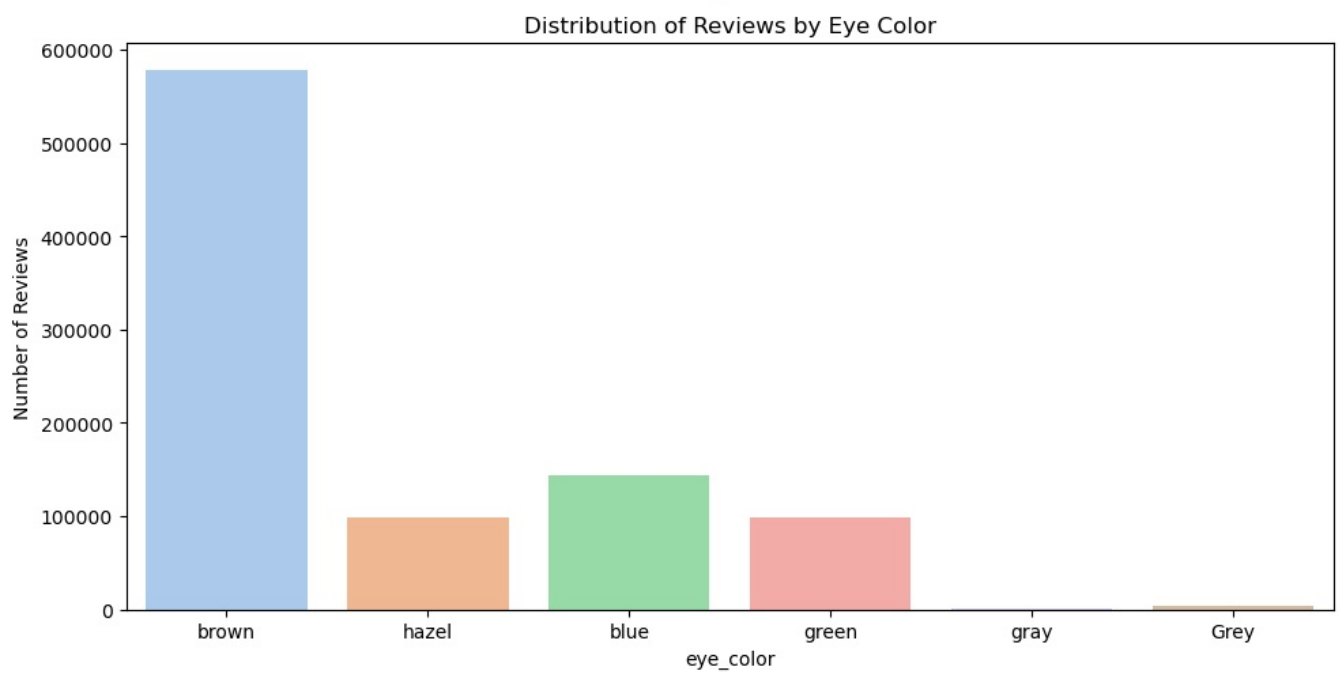
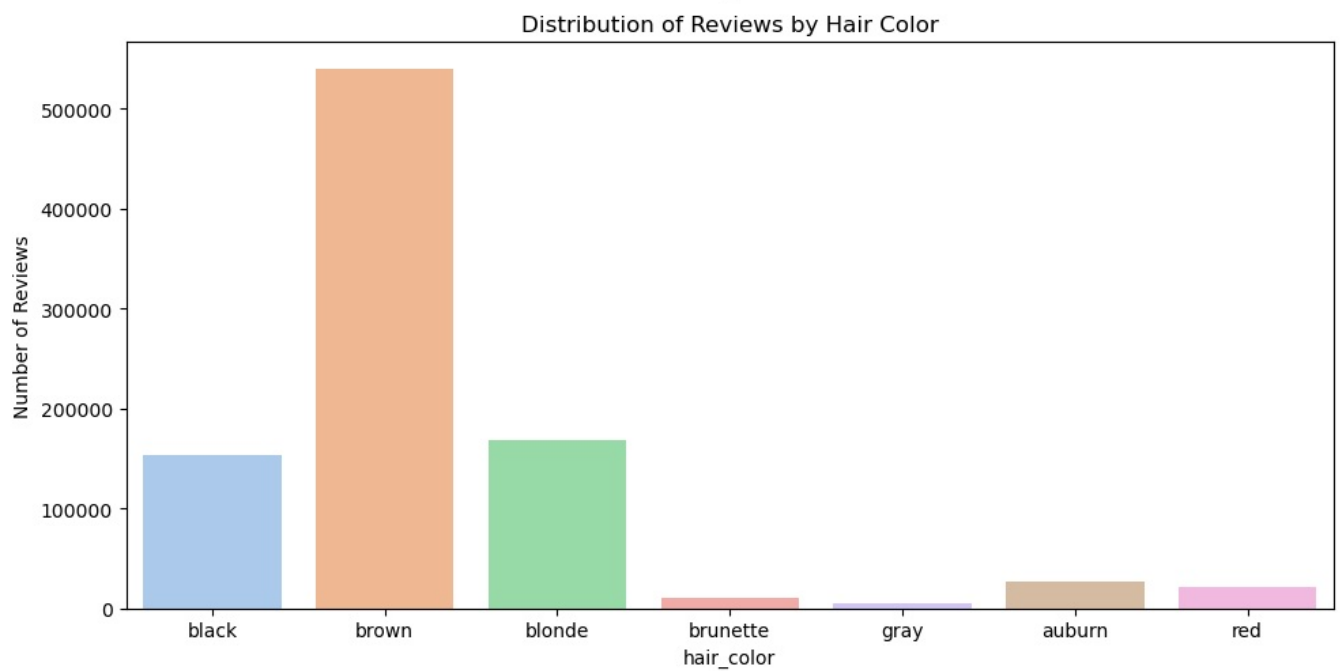
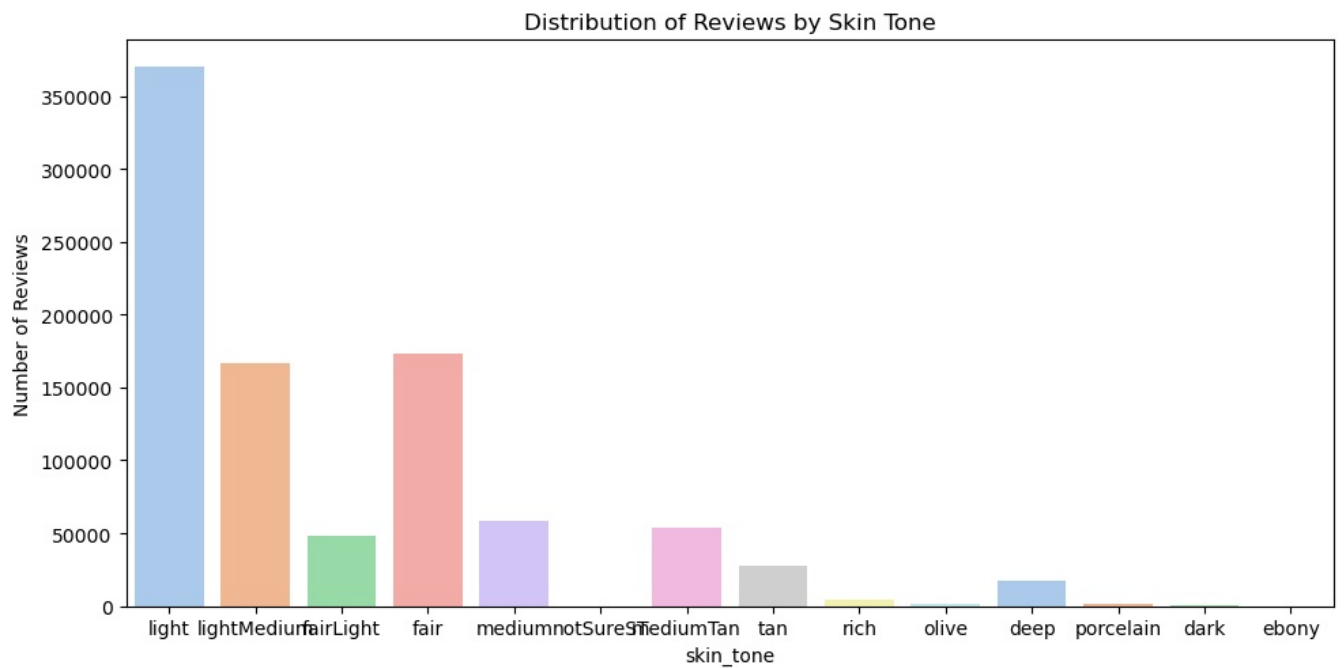
sns.barplot(x='skin_tone', y='rating', data=data, ax=axes[0], palette='pastel')
axes[0].set_title('Average Rating by Skin Tone')
axes[0].set_ylabel('Average Rating')




sns.barplot(x='hair_color', y='rating', data=data, ax=axes[1], palette='pastel')
axes[1].set_title('Average Rating by Hair Color')
axes[1].set_ylabel('Average Rating')

sns.barplot(x='eye_color', y='rating', data=data, ax=axes[2], palette='pastel')
axes[2].set_title('Average Rating by Eye Color')
axes[2].set_ylabel('Average Rating')

plt.tight_layout()
plt.show()
```







## INTERPRETATION:

### Reviews by Skin Tone:

Most reviews come from individuals with "light" and "medium" skin tones, with fewer reviews from other tones. Reviews by Hair Color: A majority of the feedback is from users with "brown" hair, followed by those with "black" and "blonde" hair.

### Reviews by Eye Color:

The majority of reviewers have "brown" eyes, followed by "blue" and "green."

### Average Rating by Skin Tone:

All skin tones generally give consistent ratings, suggesting overall satisfaction across the board.

### Average Rating by Hair and Eye Color:

Similarly, there is a consistent average rating across different hair and eye colors, indicating that products are perceived similarly across these categories.

## Solution to the Business Problem:

### Diversify Feedback Pool:

Enhance marketing campaigns targeting underrepresented skin, hair, and eye colors. Collaborations with influencers from these categories can help in acquiring diverse feedback.

### Tailored Product Lines:

Consider introducing or refining products tailored to the unique needs and preferences of underrepresented categories. For instance, certain makeup colors might be more vibrant on specific skin tones or work best for a particular hair color.

### Educational Initiatives:

Launch tutorials, guides, or workshops emphasizing the use of products for diverse skin, hair, and eye colors. This can help consumers make informed decisions and feel more represented.

### Feedback Analysis:

Conduct deeper analyses of reviews, especially from underrepresented categories, to identify any product shortcomings and areas for improvement.

By addressing the specific needs of each category and ensuring inclusivity, businesses can cater to a wider audience and improve overall customer satisfaction.

## Business Problem 3

"Are higher-priced products generally rated better than lower-priced products? We want to understand if pricing correlates with perceived product quality based on reviews."


```
In [110]: import matplotlib.pyplot as plt
import seaborn as sns

# Calculate average rating for each product
product_avg_ratings = data.groupby('product_id')['rating'].mean()

# Calculate average price for each product (it should be the same across all reviews, but taking an average for
product_avg_prices = data.groupby('product_id')['price_usd'].mean()

# Merge the two series into a DataFrame
product_info = pd.DataFrame({'avg_rating': product_avg_ratings, 'avg_price': product_avg_prices}).reset_index()

# Scatter plot for price vs average rating
plt.figure(figsize=(12, 8))
sns.scatterplot(x='avg_price', y='avg_rating', data=product_info, alpha=0.6, color='dodgerblue')
plt.title('Average Rating vs. Product Price')
plt.xlabel('Average Price (USD)')
plt.ylabel('Average Rating')
plt.grid(True, which='both', linestyle='--', linewidth=0.5)
plt.show()
```



#### INTERPRETATION:

From the scatter plot, most products, regardless of price, tend to cluster around the 4 to 5 rating mark, indicating generally positive reviews. However, there's a noticeable density of products priced under 100 *with this favorable rating*. The few products priced above 250 show a spread in ratings from 3.5 to 5. Notably, there are only a few products in the higher price range, and their ratings vary, making it challenging to deduce a clear trend for the more expensive items based solely on this data.

### Solution to the Business Problem:

#### Products Priced Under \$100:

These products generally have favorable reviews. It might be a good strategy to highlight these as 'Value for Money' or 'Budget-Friendly Picks' in marketing campaigns, emphasizing both their affordability and quality.

#### High-Priced Products (Above \$250):

Due to the variability in their ratings, it's crucial to dive deeper into individual product reviews and feedback. Identifying areas of improvement can enhance product quality and, subsequently, their ratings. Alternatively, if these products are niche or specialized, consider targeted marketing to the appropriate audience who sees the value in these items.

#### General Strategy:

It might be beneficial to conduct surveys or focus groups to better understand customer expectations from higher-priced items and if they feel the quality justifies the price. This can provide insights into whether customers equate price with quality and where the company might need to make product or pricing adjustments.

In conclusion, while there is a positive reception across all price points, the perception of value might differ. The company should prioritize enhancing the quality and value proposition of its higher-priced items and effectively market its budget-friendly yet well-reviewed products.

### Business Problem 4:


Which products have the highest counts of negative feedback? The aim is to pinpoint products that might require quality enhancements or more targeted marketing strategies to address the specific needs of different customer segments.

```
In [131]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Group by product name and sum the total negative feedback
negative_feedback_summary = text_edited.groupby('product_name')['total_neg_feedback_count'].sum().sort_values(ascending=True)

# Extract the top 10 products with the highest negative feedback
top_negative_products = negative_feedback_summary.head(10)

# Visualization
plt.figure(figsize=(12, 8))
sns.barplot(y=top_negative_products.index, x=top_negative_products.values, palette='Reds_r')
plt.title('Top 10 Products with Highest Negative Feedback')
plt.xlabel('Total Negative Feedback Count')
plt.ylabel('Product Name')
plt.show()
```



## INTERPRETATION:

The graph illustrates the top 10 products based on the volume of negative feedback they received. At the top, "AHA 30% + BHA 2% Exfoliating Peeling Solution" stands out with the highest negative feedback, followed by "Niacinamide 10% + Zinc 1% Oil Control Serum" and "Lip Sleeping Mask Intense Hydration with Vitamin C". This suggests that customers may have encountered issues or unsatisfactory results with these products more frequently than others on the list.

## Solution to the Business Problem:

### Quality Assessment & Product Refinement:

Start with a thorough investigation into why these specific products have high negative feedback. Analyze the common themes in the negative feedback and assess if there are consistent issues being reported. For example, is the product causing skin irritation, or are customers finding it ineffective for its intended purpose? Use this feedback to make necessary formula adjustments or address other product-related concerns.

### Engage in Direct Communication:

Reach out to customers who provided negative feedback to understand their concerns better and potentially offer solutions or compensation. This can not only help improve the product but also rebuild trust with dissatisfied customers.

### Educational Marketing:

Perhaps customers are using the products incorrectly or have misunderstood the product's intended purpose. In such cases, launching educational marketing campaigns, tutorials, or how-to guides can help customers achieve desired results with the product.

### Targeted Marketing & Promotions:

If a certain customer segment is particularly dissatisfied, consider targeted marketing strategies that address their specific needs and concerns. Additionally, offering promotions or discounts on these products can incentivize users to give them another try, especially if improvements have been made.

### Product Demonstrations & Sampling:

Organize events or collaborations with beauty experts to demonstrate the correct usage of these products. Offering samples can also provide potential customers with a risk-free way to try the product before committing to a purchase.

By addressing the concerns directly and ensuring products meet customers' needs and expectations, the company can potentially reduce negative feedback and enhance overall brand perception.

## Business Problem 5:

In today's digital era, where customer reviews and recommendations greatly influence purchasing behavior, understanding and anticipating customer sentiment is crucial. For a global beauty and skincare retailer like Sephora, every product recommendation, or lack thereof, can significantly impact sales, brand perception, and customer loyalty.

At the heart of this challenge lies a critical question: Can we predict if a customer will recommend a product based on their interactions, ratings, and feedback? And if so, how can such predictions shape Sephora's strategic decisions, from product placements, inventory management, to marketing campaigns?

```
In [111.. # Calculate class distribution
class_distribution = data['is_recommended'].value_counts()

# Calculate class proportions
class_proportions = class_distribution / class_distribution.sum()

print("Class Proportions:")
print(class_proportions)

Class Proportions:
1.0    0.86564
0.0    0.13436
Name: is_recommended, dtype: float64
```

```
In [112.. pip install imbalanced-learn

Requirement already satisfied: imbalanced-learn in c:\users\logeshwar\anaconda3\lib\site-packages (0.11.0)Note:
you may need to restart the kernel to use updated packages.

Requirement already satisfied: joblib>=1.1.1 in c:\users\logeshwar\anaconda3\lib\site-packages (from imbalanced
-learn) (1.3.2)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\logeshwar\anaconda3\lib\site-packages (from imba
lanced-learn) (1.0.2)
Requirement already satisfied: scipy>=1.5.0 in c:\users\logeshwar\anaconda3\lib\site-packages (from imbalanced-
learn) (1.9.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\logeshwar\anaconda3\lib\site-packages (from imb
alanced-learn) (2.2.0)
Requirement already satisfied: numpy>=1.17.3 in c:\users\logeshwar\anaconda3\lib\site-packages (from imbalanced
-learn) (1.21.5)
```

```
In [113.. data['submission_time'] = pd.to_datetime(data['submission_time'])
```

```
In [114.. text_edited.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1301136 entries, 0 to 49976
Data columns (total 17 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   author_id             1301136 non-null object
 1   rating                1301136 non-null int64
 2   is_recommended        1301136 non-null object
 3   total_feedback_count  1301136 non-null int64
 4   total_neg_feedback_count 1301136 non-null int64
 5   total_pos_feedback_count 1301136 non-null int64
 6   submission_time       1301136 non-null datetime64[ns]
 7   review_text           1301136 non-null object
 8   review_title          1301136 non-null object
 9   skin_tone              1301136 non-null object
10   eye_color              1301136 non-null object
11   skin_type              1301136 non-null object
12   hair_color             1301136 non-null object
13   product_id            1301136 non-null object
14   product_name           1301136 non-null object
15   brand_name             1301136 non-null object
16   price_usd              1301136 non-null float64
dtypes: datetime64[ns](1), float64(1), int64(4), object(11)
memory usage: 178.7+ MB
```

```
In [115.. correlation = data['rating'].corr(data['price_usd'])
print("Correlation between rating and price_usd:", correlation)

Correlation between rating and price_usd: -0.0026157697373409665
```

The correlation value of  $-0.0026$  between rating and price\_usd is very close to 0. This indicates that there's a negligible linear relationship between the two variables in the dataset.

In practical terms, this suggests that the price\_usd of a product has almost no linear effect on its rating. The rating of a product is not

increasing or decreasing significantly with its price.

# Machine Learning Model

## Classification Model - Predicting Product Recommendation:

"Enhancing Brand Trust and Customer Loyalty through Predictive Analysis of Product Recommendations at Sephora"

```
In [119]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Load the dataset
data = pd.read_csv(r"C:\Users\logeshwar\Downloads\Sephora Reviews Dataset.csv") # Replace with the actual path

# Selecting features for classification
features_classification = ['rating', 'total_feedback_count', 'total_neg_feedback_count', 'total_pos_feedback_co
X_classification = data[features_classification]
y_classification = data['is_recommended']

# Splitting data into training and testing sets
X_train_classification, X_test_classification, y_train_classification, y_test_classification = train_test_split

# Building the Random Forest Classifier model
model_classification = RandomForestClassifier(random_state=42)
model_classification.fit(X_train_classification, y_train_classification)

# Making predictions on the test set
y_pred_classification = model_classification.predict(X_test_classification)

# Evaluating the classification model
accuracy_classification = accuracy_score(y_test_classification, y_pred_classification)
print("Classification Model - Accuracy:", accuracy_classification)
print(classification_report(y_test_classification, y_pred_classification))
```

```
Classification Model - Accuracy: 0.9452927035280358
      precision    recall  f1-score   support

      0.0         0.77         0.84         0.81         24837
      1.0         0.98         0.96         0.97         159855

 accuracy         0.95         0.95         0.95         184692
 macro avg        0.87         0.90         0.89         184692
 weighted avg     0.95         0.95         0.95         184692
```

```
In [120]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Load the dataset
data = pd.read_csv(r"C:\Users\logeshwar\Downloads\Sephora Reviews Dataset.csv") # Replace with the actual path

# Selecting features for classification
features_classification = ['rating', 'total_feedback_count', 'total_neg_feedback_count', 'total_pos_feedback_co
X_classification = data[features_classification]
y_classification = data['is_recommended']

# Splitting data into training and testing sets
X_train_classification, X_test_classification, y_train_classification, y_test_classification = train_test_split

# Building the Random Forest Classifier model
model_classification = RandomForestClassifier(random_state=42)
model_classification.fit(X_train_classification, y_train_classification)

# Making predictions on the test set
y_pred_classification = model_classification.predict(X_test_classification)

# Evaluating the classification model
accuracy_classification = accuracy_score(y_test_classification, y_pred_classification)
print("Classification Model - Accuracy:", accuracy_classification)
print(classification_report(y_test_classification, y_pred_classification))

# Interpreting the Results
# Analyze the classification report and feature importance

# Hyperparameter Tuning
# You can use techniques like grid search or randomized search to find optimal hyperparameters

# Feature Importance
feature_importance = model_classification.feature_importances_
```

```
print("Feature Importance:", feature_importance)

# Model Deployment (if desired)
# Deploy the model in a practical application if it meets your requirements

# Monitoring and Maintenance (if deployed)
# Monitor the model's performance and retrain as needed

# Iterative Improvement (if necessary)
# Refine the model by iterating on data preprocessing, feature selection, or trying different algorithms
```

```
Classification Model - Accuracy: 0.9452927035280358
      precision    recall  f1-score   support
```

```

0.0         0.77         0.84         0.81         24837
1.0         0.98         0.96         0.97         159855

 accuracy
macro avg          0.87         0.90         0.89         184692
weighted avg       0.95         0.95         0.95         184692
```

```
Feature Importance: [0.8016889  0.03336477 0.04747366 0.02108383 0.09638885]
```

## INTERPRETATIONS:

In summary, while our model has a high overall accuracy, it's important to pay attention to its performance on both classes, especially the minority class. Depending on our specific goals and the consequences of misclassification, you may need to take additional steps to address the class imbalance and improve the model's ability to correctly classify both recommended and not recommended reviews.

# Resampling Techniques

Handling Imbalanced Data:

Using SMOTE for oversampling:

```
In [121] from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_classification, y_train_classification)
```

## INTERPRETATIONS:

SMOTE to tackle the problem of having too many of one class and too few of another in your training data.

SMOTE creates new, synthetic samples of the minority class (the class with fewer samples) to balance the dataset.

By using SMOTE:

We're trying to make your model perform better for both classes by having equal samples of each class in the training data. We've only applied it to the training data, so your test data remains untouched. While this can improve performance, sometimes it can also add noise since the new samples are artificially created. In short: We're using SMOTE to balance your classes in the training set to help your model predict better.

```
In [122] text_edited.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1301136 entries, 0 to 49976
Data columns (total 17 columns):
 #   Column                                Non-Null Count  Dtype  
---  -
 0   author_id                            1301136 non-null object  
 1   rating                               1301136 non-null int64   
 2   is_recommended                       1301136 non-null object  
 3   total_feedback_count                 1301136 non-null int64   
 4   total_neg_feedback_count             1301136 non-null int64   
 5   total_pos_feedback_count             1301136 non-null int64   
 6   submission_time                      1301136 non-null datetime64[ns]
 7   review_text                          1301136 non-null object  
 8   review_title                         1301136 non-null object  
 9   skin_tone                            1301136 non-null object  
10   eye_color                            1301136 non-null object  
11   skin_type                            1301136 non-null object  
12   hair_color                           1301136 non-null object  
13   product_id                           1301136 non-null object  
14   product_name                         1301136 non-null object  
15   brand_name                           1301136 non-null object  
16   price_usd                            1301136 non-null float64  
dtypes: datetime64[ns](1), float64(1), int64(4), object(11)
memory usage: 178.7+ MB
```

```
In [123] importances = model_classification.feature_importances_
```



```
In [124.. from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Load the dataset
data = pd.read_csv(r"C:\Users\logeshwar\Downloads\Sephora Reviews Dataset.csv") # Replace with the actual path

# Selecting features for classification
features_classification = ['rating', 'total_feedback_count', 'total_neg_feedback_count', 'total_pos_feedback_co
X_classification = data[features_classification]
y_classification = data['is_recommended']

# Splitting data into training and testing sets
X_train_classification, X_test_classification, y_train_classification, y_test_classification = train_test_split

# Building the Random Forest Classifier model
model_classification = RandomForestClassifier(random_state=42)
model_classification.fit(X_train_classification, y_train_classification)

# Making predictions on the test set
y_pred_classification = model_classification.predict(X_test_classification)

# Evaluating the classification model
accuracy_classification = accuracy_score(y_test_classification, y_pred_classification)
print("Classification Model - Accuracy:", accuracy_classification)
print(classification_report(y_test_classification, y_pred_classification))

# Interpreting the Results
# Analyze the classification report and feature importance

# Hyperparameter Tuning
# You can use techniques like grid search or randomized search to find optimal hyperparameters

# Feature Importance
feature_importance = model_classification.feature_importances_
print("Feature Importance:", feature_importance)

# Model Deployment (if desired)
# Deploy the model in a practical application if it meets your requirements

# Monitoring and Maintenance (if deployed)
# Monitor the model's performance and retrain as needed

# Iterative Improvement (if necessary)
# Refine the model by iterating on data preprocessing, feature selection, or trying different algorithms
```

Classification Model - Accuracy: 0.9452927035280358

	precision	recall	f1-score	support
0.0	0.77	0.84	0.81	24837
1.0	0.98	0.96	0.97	159855
accuracy			0.95	184692
macro avg	0.87	0.90	0.89	184692
weighted avg	0.95	0.95	0.95	184692

Feature Importance: [0.8016889 0.03336477 0.04747366 0.02108383 0.09638885]

```
In [125.. pip install xgboost
```

Requirement already satisfied: xgboost in c:\users\logeshwar\anaconda3\lib\site-packages (2.0.0)Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: scipy in c:\users\logeshwar\anaconda3\lib\site-packages (from xgboost) (1.9.1)

Requirement already satisfied: numpy in c:\users\logeshwar\anaconda3\lib\site-packages (from xgboost) (1.21.5)

```
In [126.. import pandas as pd
from sklearn.model_selection import train_test_split
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, classification_report

# Load the dataset
data = pd.read_csv(r"C:\Users\logeshwar\Downloads\Sephora Reviews Dataset.csv") # Replace with the actual path

# Selecting features for classification
features_classification = ['rating', 'total_feedback_count', 'total_neg_feedback_count', 'total_pos_feedback_co
X_classification = data[features_classification]
y_classification = data['is_recommended']

# Splitting data into training and testing sets
X_train_classification, X_test_classification, y_train_classification, y_test_classification = train_test_split

# Building the XGBoost Classifier model
model_classification = XGBClassifier(random_state=42, use_label_encoder=False, eval_metric="logloss")
model_classification.fit(X_train_classification, y_train_classification)

# Making predictions on the test set
y_pred_classification = model_classification.predict(X_test_classification)
```

```
# Evaluating the classification model
accuracy_classification = accuracy_score(y_test_classification, y_pred_classification)
print("Classification Model - Accuracy:", accuracy_classification)
print(classification_report(y_test_classification, y_pred_classification))

# Feature Importance
feature_importance = model_classification.feature_importances_
for feature, importance in zip(features_classification, feature_importance):
    print(f"Feature: {feature}, Importance: {importance}")
```

```
Classification Model - Accuracy: 0.9465705065731055
precision    recall  f1-score   support
```

```
0.0         0.76     0.88     0.82     24837
1.0         0.98     0.96     0.97     159855

accuracy
macro avg      0.87     0.92     0.89     184692
weighted avg   0.95     0.95     0.95     184692
```

```
Feature: rating, Importance: 0.9732348918914795
Feature: total_feedback_count, Importance: 0.003195783356204629
Feature: total_neg_feedback_count, Importance: 0.003949116449803114
Feature: total_pos_feedback_count, Importance: 0.006272484548389912
Feature: price_usd, Importance: 0.01334768533706665
```

```
In [127]: import pandas as pd
from sklearn.model_selection import train_test_split
from xgboost import XGBClassifier
from sklearn.calibration import CalibratedClassifierCV
from sklearn.metrics import accuracy_score, classification_report

# Load the dataset
data = pd.read_csv(r"C:\Users\logeshwar\Downloads\Sephora Reviews Dataset.csv")

# Selecting features for classification
features_classification = ['rating', 'total_feedback_count', 'total_neg_feedback_count', 'total_pos_feedback_count']
X_classification = data[features_classification]
y_classification = data['is_recommended']

# Splitting data into training and testing sets
X_train_classification, X_test_classification, y_train_classification, y_test_classification = train_test_split(X_classification, y_classification)

# Building the XGBoost Classifier model
xgb_model = XGBClassifier(random_state=42, use_label_encoder=False, eval_metric="logloss")
xgb_model.fit(X_train_classification, y_train_classification)

# Calibrating the model
calibrated = CalibratedClassifierCV(xgb_model, method='sigmoid', cv='prefit')
calibrated.fit(X_train_classification, y_train_classification)

# Making predictions on the test set
y_pred_classification = calibrated.predict(X_test_classification)

# Evaluating the calibrated classification model
accuracy_classification = accuracy_score(y_test_classification, y_pred_classification)
print("Calibrated Model - Accuracy:", accuracy_classification)
print(classification_report(y_test_classification, y_pred_classification))
```

```
Calibrated Model - Accuracy: 0.9455417668334308
precision    recall  f1-score   support
```

```
0.0         0.79     0.81     0.80     24837
1.0         0.97     0.97     0.97     159855

accuracy
macro avg      0.88     0.89     0.88     184692
weighted avg   0.95     0.95     0.95     184692
```

## Overall Interpretation:

The model performs exceptionally well in predicting when a product will be recommended, with both precision and recall values at 97%. However, for products that aren't recommended, the model's performance, while decent, isn't as stellar, with precision and recall values around 79% and 81% respectively. Given the significant class imbalance (much more "Recommended" than "Not Recommended"), the weighted average scores are strongly influenced by the "Recommended" class, leading to an overall high accuracy. This suggests the model is more reliable when predicting positive recommendations, but there's still room for improvement when identifying negative ones.

## Solution to the business problem:

## Background:

Sephora is a leading beauty and skincare retailer. As with many e-commerce platforms, customer reviews play a pivotal role in influencing purchasing decisions. Reviews, particularly whether a product is recommended or not, can significantly affect sales and brand trust.

## Objective:

To predict whether a product will be recommended by a customer based on various factors like rating, feedback counts, and product price. By predicting this, Sephora aims to:

### Improve Customer Experience:

By understanding the factors that lead to product recommendations, Sephora can prioritize showcasing products that are more likely to be recommended, leading to higher customer satisfaction.

### Increase Sales:

Products with higher recommendation rates tend to be purchased more. By predicting and subsequently highlighting such products, there's potential to drive more sales.

### Enhance Brand Trust:

A product that has a higher rate of recommendations is perceived as trustworthy. By ensuring that these products are prominently featured, Sephora can increase trust among its customers.

### Product & Inventory Management:

Products that aren't likely to be recommended can be re-evaluated, and inventory decisions can be optimized.

### Tailored Marketing:

Products predicted to have high recommendation rates can be featured in marketing campaigns, emails, and promotions.

## Model Use:

The classification model, built using XGBoost and subsequently calibrated, predicts if a customer would recommend a product based on selected features. With an accuracy of approximately 94.5%, the model does an excellent job in making this prediction. This high accuracy means the model's predictions are reliable and can be used for strategic decisions.

## Evaluation:

The overall accuracy is approximately 94.5%, suggesting the model does well in classifying the reviews. The precision, recall, and F1-score for the "recommended" class (1.0) are notably high. This is crucial since this class represents positive recommendations, which are vital for business. The metrics for the "not recommended" class (0.0) are lower compared to the "recommended" class but are still decent, showing the model's capability in identifying negative reviews as well.

## Future Steps:

Investigate the reasons behind the products that are not recommended and find ways to address the underlying issues. Use the model to dynamically adjust the display of products on the website based on their predicted recommendation status. Periodically retrain the model with new reviews to keep it updated and accurate. This approach, rooted in data-driven decision-making, can provide Sephora with a competitive edge in the increasingly crowded online beauty and skincare market.